

CS 61C: Great Ideas in Computer Architecture

Lecture 24: More I/O: DMA, Disks, Networking

Instructor: Sagar Karandikar
sagark@eecs.berkeley.edu

<http://inst.eecs.berkeley.edu/~cs61c>
Berkeley EECS
ELECTRICAL ENGINEERING & COMPUTER SCIENCES

Review: I/O

- Programmed I/O:
 - CPU execs lw/sw instructions for all data movement to/from devices
 - CPU spends time doing 3 things:
 - Getting data from device to main mem.
 - Using data to compute
 - Sending data back to main mem.
- Polling
- Interrupts

2

Working with real devices

- Programmed I/O: DMA
 - CPU execs lw/sw instructions for all data movement to/from devices
 - CPU spends time doing 3 things:
 - Getting data from device to main mem.
 - Using data to compute
 - Sending data back to main mem.
- Polling
- Interrupts

3

Agenda

- Direct Memory Access (DMA)
- Disks
- Networking

4

What's wrong with Programmed I/O?

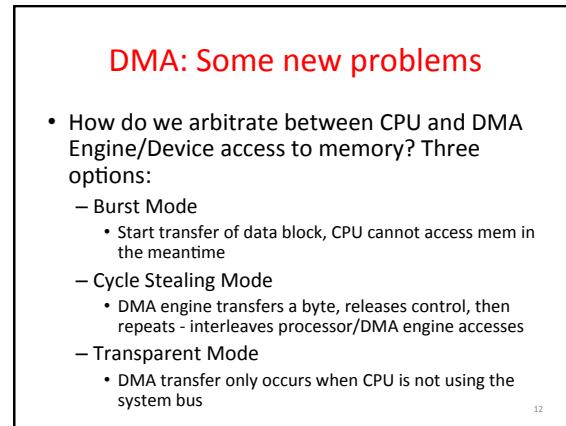
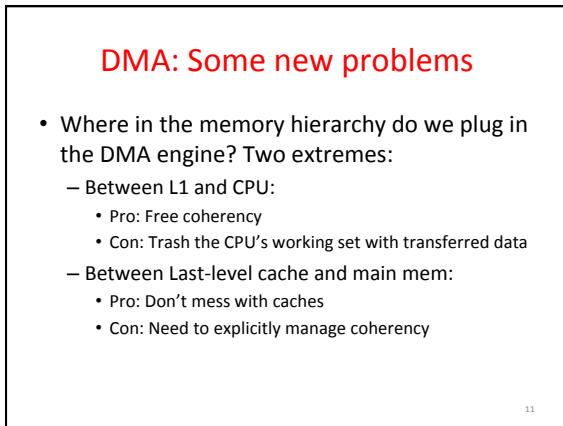
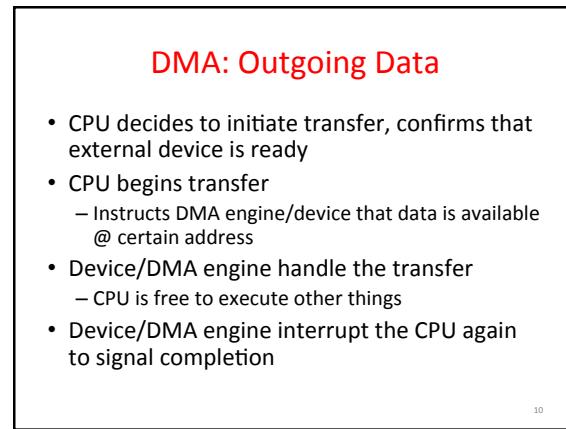
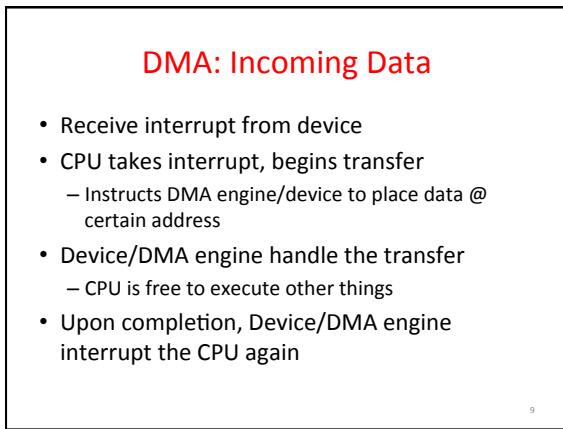
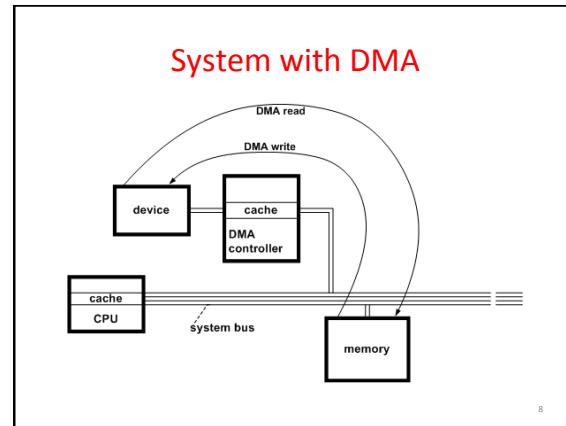
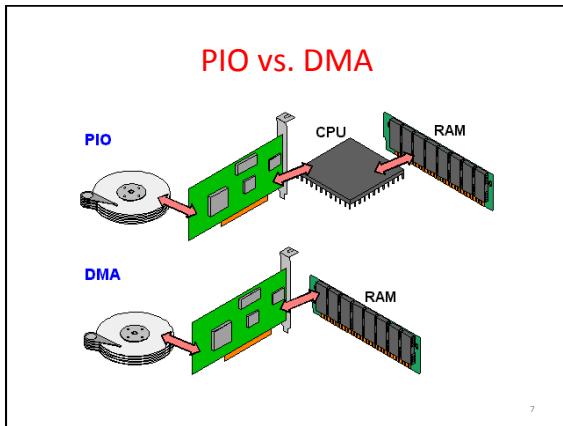
- CPU has sole control over main memory
- Not ideal since...
 - CPU has to execute all transfers, could be doing other work
 - Device speeds don't align well with CPU speeds
 - Energy cost of using beefy general-purpose CPU where simpler hardware would suffice

5

Direct Memory Access (DMA)

- Allows other devices to directly read/write from main memory
- New Hardware: the DMA Engine
- DMA engine contains registers written by CPU:
 - Memory address to place data
 - # of bytes
 - I/O port #, direction of transfer
 - unit of transfer, amount to transfer per burst

6



Administrivia

- Project 4 out now
 - Competition for glory + EC, awards during last lecture
- HW6 Out – VM, I/O, Parity, ECC
- Guerrilla Section on VM, I/O, ECC, on Thursday from 5-7pm, Woz

13

Administrivia

- Upcoming Lecture Schedule
 - 8/04: I/O: DMA, Disks, Networking (today)
 - 8/05: Dependability: Parity, ECC, RAID
 - Last day of new material
 - 8/06: Final Exam Review, Fa14 Final
 - 8/10: Final Exam Review, Sp15 Final
 - 8/11: Summary, What's Next? (+ HKN reviews)
 - Project 4 Competition Winners Announced
 - 8/12: No Lecture, I'll have OH in this room

14

Administrivia

- Final Exam is next Thursday (8/13)
 - 9am-12pm, 10 Evans
 - Three 8.5"x11", double-sided, handwritten cheat-sheets
 - Broken into MT1 section, MT2 section, and post-MT2 section (so that the clobber can be applied)
 - TA Review Session: 08/09, 1-5pm, HP Aud.
 - In-class review: 08/06, 08/10

15

Agenda

- Direct Memory Access (DMA)
- **Disks**
- Networking

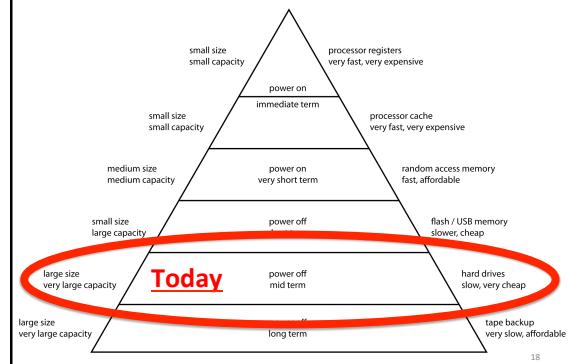
16

Review - 6 Great Ideas in Computer Architecture

1. Layers of Representation/Interpretation
2. Moore's Law
- 3. Principle of Locality/Memory Hierarchy**
4. Parallelism
5. Performance Measurement & Improvement
6. Dependability via Redundancy

17

Computer Memory Hierarchy



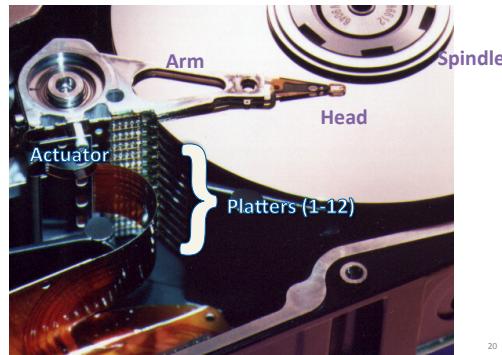
18

Magnetic Disk – common I/O device

- A kind of computer memory
 - Information stored by magnetizing ferrite material on surface of rotating disk
 - similar to tape recorder except digital rather than analog data
- A type of non-volatile storage
 - retains its value without applying power to disk.
- Two Types of Magnetic Disk
 - Floppy disks – slower, less dense, removable.
 - Hard Disk Drives (HDD) – faster, more dense, non-removable.
- Purpose in computer systems (Hard Drive):
 - Long-term, inexpensive storage for files
 - “Backup” for main-memory. Large, inexpensive, slow level in the memory hierarchy (virtual memory)

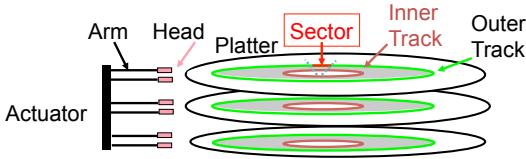
19

Photo of Disk Head, Arm, Actuator



20

Disk Device Terminology



- Several platters, with information recorded magnetically on both surfaces (usually)
- Bits recorded in **tracks**, which in turn divided into **sectors** (e.g., 512 Bytes)
- Actuator** moves **head** (end of **arm**) over track ("seek"), wait for **sector** rotate under **head**, then read or write

21

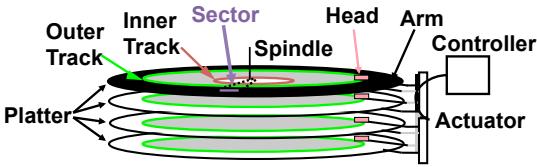
Hard Drives are Sealed. Why?

- The closer the head to the disk, the smaller the “spot size” and thus the denser the recording.
 - Measured in Gbit/in²
 - ~900 Gbit/in² is state of the art
 - Started out at 2 Kbit/in²
 - ~450,000,000x improvement in ~60 years
- Disks are sealed to keep the dust out.
 - Heads are designed to “fly” at around 5-20nm above the surface of the disk.
 - 99.999% of the head/arm weight is supported by the air bearing force (air cushion) developed between the disk and the head.



22

Disk Device Performance (1/2)



- Disk Access Time = Seek Time + Rotation Time + Transfer Time + Controller Overhead**
 - Seek Time = time to position the head assembly at the proper cylinder
 - Rotation Time = time for the disk to rotate to the point where the first sectors of the block to access reach the head
 - Transfer Time = time taken by the sectors of the block and any gaps between them to rotate past the head

23

Disk Device Performance (2/2)

- Average values to plug into the formula:
- Rotation Time: Average distance of sector from head?
 - 1/2 time of a rotation
 - 7200 Revolutions Per Minute \Rightarrow 120 Rev/sec
 - 1 revolution $= 1/120$ sec \Rightarrow 8.33 milliseconds
 - 1/2 rotation (revolution) \Rightarrow 4.17 ms
- Seek time: Average no. tracks to move arm?
 - Number of tracks/3 (see CS186 for the math)
 - Then, Seek time = number of tracks moved * time to move across one track

24

But wait!

- Performance estimates are different in practice
- Lots of “magic” going on behind the scenes in disks. One example:
 - Many disks have on-disk caches, which are completely hidden from the outside world
 - Previous formula completely replaced with on-disk cache access time

25

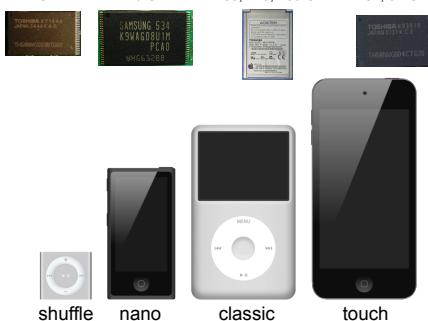
Where does Flash memory come in?

- ~5-10 years ago: Microdrives and Flash memory (e.g., CompactFlash) went head-to-head
 - Both non-volatile (no power, data ok)
 - Flash benefits: durable & lower power (no moving parts, need to spin µdrives up/down)
 - Disk cost = fixed cost of motor + arm mechanics, but actual magnetic media cost very low
 - Flash cost = most cost/bit of flash chips
 - Over time, cost/bit of flash came down, became cost competitive

26

What does Apple put in its iPods?

Toshiba flash 2 GB Samsung flash 16 GB Toshiba 1.8-inch HDD 80, 120, 160 GB Toshiba flash 32, 64 GB



27

What does Apple put in its iPods?

Toshiba flash 2 GB Samsung flash 16 GB Toshiba 1.8-inch HDD 80, 120, 160 GB Toshiba flash 32, 64 GB



28

Flash Memory / SSDs

- How does Flash memory work?
 - NMOS transistor with an additional conductor between gate and source/drain which “traps” electrons. The presence/absence is a 1 or 0
- Requires complex management algorithms to avoid wearing out cells
 - Early SSDs had severe reliability issues due to poorly implemented firmware
 - More in CS162



29

iClicker Question

- We have the following disk:
 - 15000 Cylinders, 1 ms to cross 1000 Cylinders
 - 15000 RPM = 4 ms per rotation
 - Want to copy 1 MB, transfer rate of 1000 MB/s
 - 1 ms controller processing time
- What is the access time using our model?

$$\text{Disk Access Time} = \text{Seek Time} + \text{Rotation Time} + \text{Transfer Time} + \text{Controller Processing Time}$$

A	B	C	D	E
10.5 ms	9 ms	8.5 ms	11.4 ms	12 ms

30

iClicker Question

- We have the following disk:

- 15000 Cylinders, 1 ms to cross 1000 Cylinders
- 15000 RPM = 4 ms per rotation
- Want to copy 1 MB, transfer rate of 1000 MB/s
- 1 ms controller processing time

- What is the access time?

$$\text{Seek} = \# \text{ cylinders}/3 * \text{time} = 15000/3 * 1\text{ms}/1000 \text{ cylinders} = 5\text{ms}$$

$$\text{Rotation} = \text{time for } \frac{1}{2} \text{ rotation} = 4 \text{ ms} / 2 = 2 \text{ ms}$$

$$\text{Transfer} = \text{Size} / \text{transfer rate} = 1 \text{ MB} / (1000 \text{ MB/s}) = 1 \text{ ms}$$

$$\text{Controller} = 1 \text{ ms}$$

$$\text{Total} = 5 + 2 + 1 + 1 = 9 \text{ ms}$$

31

<http://research.microsoft.com/apps/pubs/default.aspx?id=212001>

CS61C (and CS150) in the News

Microsoft “Catapult”, ISCA 2014

FPGAs are “programmable” hardware used by computer architects and digital circuit designers, lie somewhere between CPUs and custom chips (ASICs).

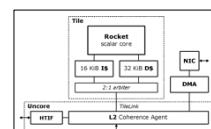


“Microsoft published a paper at ISCA about using FPGAs in datacenters for page ranking processing for Bing. In a test deployment, MS reported up to 95% more throughput for only 10% more power. The added TCO was less than 30%. Microsoft used Altera Stratix V FPGAs in a PCIe form-factor with 8GB of DDR3 RAM on each board. The FPGAs were connected using a 10Gb SAS network.” - AnandTech



Hardware Acceleration of Key-Value Stores

- Datacenter apps, path through CPU/kernel/app ≈ 86% of request latency



Base System: First Physical RISC-V System with Networking Support

Sagar Karandikar

Howard Mao

Albert Ou

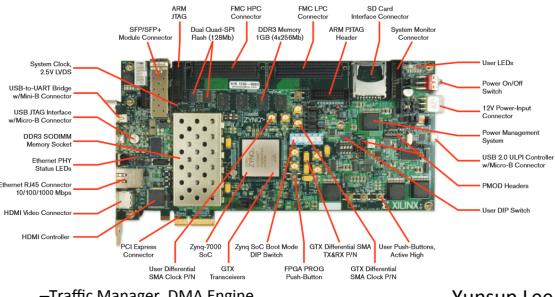
Yunsup Lee

Krste Asanovic

- Goal: Serve popular Key-Value Store GET requests without CPU
- Soft-managed cache attached to NIC, RoCC CPU interface
- Benchmarking on FPGA:
 - RISC-V Rocket @ 50 MHz
 - NIC from TEMAC/PCS-PMA + 1 Gb SFP
 - Hardware KV-Store Accelerator
 - Traffic Manager, DMA Engine
- Written in Chisel



Hardware Acceleration of Key-Value Stores



–Traffic Manager, DMA Engine

• Written in Chisel

Yunsup Lee

Krste Asanovic

Break

35

Agenda

- Direct Memory Access (DMA)
- Disks
- Networking**

36

Networks: Talking to the Outside World

- Originally sharing I/O devices between computers
 - E.g., printers
- Then communicating between computers
 - E.g., file transfer protocol
- Then communicating between people
 - E.g., e-mail
- Then communicating between networks of computers
 - E.g., file sharing, www, ...

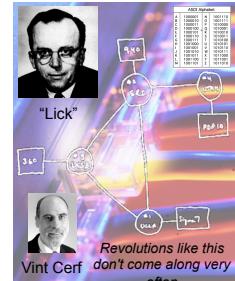
8/2/15

37

www.computerhistory.org/internet_history

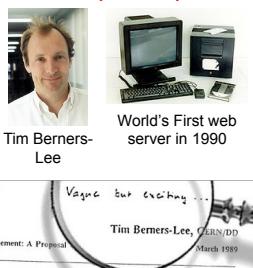
The Internet (1962)

- History
 - JCR Licklider, as head of ARPA, writes on "intergalactic network"
 - 1963 : ASCII becomes first universal computer standard
 - 1969 : Defense Advanced Research Projects Agency (DARPA) deploys 4 "nodes" @ UCLA, SRI, Utah, & UCSB
 - 1973 Robert Kahn & Vint Cerf invent TCP, now part of the Internet Protocol Suite
- Internet growth rates
 - Exponential since start!

8/2/15 en.wikipedia.org/wiki/Internet_Protocol_Suite 38en.wikipedia.org/wiki/History_of_the_World_Wide_Web

The World Wide Web (1989)

- "System of interlinked hypertext documents on the Internet"
- History
 - 1945: Vannevar Bush describes hypertext system called "memex" in article
 - 1989: Sir Tim Berners-Lee proposes, gets system up '90
 - ~2000 Dot-com entrepreneurs rushed in, 2001 bubble burst
- Wayback Machine
 - Snapshots of web over time
- Today : Access anywhere!

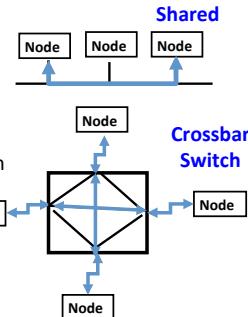


8/2/15

39

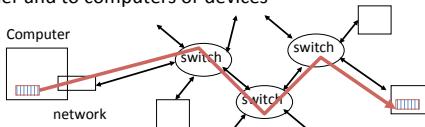
Shared vs. Switch-Based Networks

- Shared vs. Switched:
 - **Shared:** 1 at a time (CSMA/CD)
 - **Switched:** pairs ("point-to-point" connections) communicate at same time
- Aggregate bandwidth (BW) in switched network is many times shared:
 - point-to-point faster since no arbitration, simpler interface



40

What makes networks work?

- links connecting switches and/or routers to each other and to computers or devices
 
- ability to name the components and to route packets of information - messages - from a source to a destination
- Layering, redundancy, protocols, and encapsulation as means of abstraction (61C big idea)

8/2/15

41

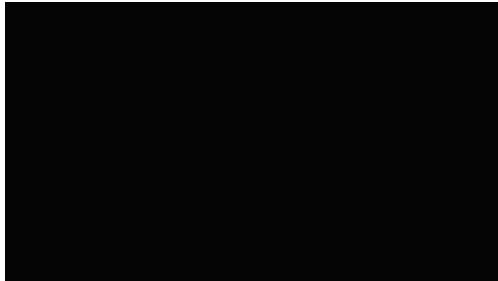
Software Protocol to Send and Receive

- SW Send steps
 - 1: Application copies data to OS buffer
 - 2: OS calculates checksum, starts timer
 - 3: OS sends data to network interface HW and says start
- SW Receive steps
 - 3: OS copies data from network interface HW to OS buffer
 - 2: OS calculates checksum, if OK, send ACK; if not, delete message (sender resends when timer expires)
 - 1: If OK, OS copies data to user address space, & signals application to continue

Dest	Src	Checksum					
Net ID	Net ID	Len	ACK INFO	CMD/ Address /Data	Header	Payload	Trailer

42

Networks are like Ogres



https://www.youtube.com/watch?v=_bMcXVe8zls

43

Protocol for Networks of Networks?

What does it take to send packets across the globe?

- Bits on wire
- Packets on wire
- Delivery packets within a single physical network
- Deliver packets across multiple networks
- Ensure the destination received the data
- Do something with the data

Protocol for Networks of Networks?

- Abstraction to cope with complexity of communication

- Networks are like **egress onions**

- **Hierarchy of layers:**

- Application (chat client, game, etc.)
- Transport (TCP, UDP)
- Network (IP)
- Data Link Layer (ethernet)
- Physical Link (copper, wireless, etc.)

Networks are like onions.
They stink?
Yes. No!
Oh, they make you cry.
Not... Layers.
Onions have layers.
Networks have layers.

45

Protocol Family Concept

- Key to **protocol families** is that communication occurs **logically** at the same level of the protocol, called **peer-to-peer**...

...but is **implemented** via services at the next **lower level**

- **Encapsulation:** carry higher level information within lower level “envelope”

8/2/15

46

Inspiration...

- CEO A writes letter to CEO B
 - Folds letter and hands it to assistant **Dear John,**
- Your days are numbered.
 - Puts letter in envelope with CEO B's full name
 - Takes to FedEx
- FedEx Office
 - Puts letter in larger envelope
 - ~~Puts part~~ and street address on FedEx envelope
 - Puts package on FedEx delivery truck
- FedEx delivers to other company

The Path of the Letter

“Peers” on each side understand the same things
No one else needs to
Lowest level has most packaging



