

CS 61C: Great Ideas in Computer Architecture

Lecture 28: *Course Summary and Wrap-Up*

Instructor: Sagar Karandikar
sagark@eecs.berkeley.edu

<http://inst.eecs.berkeley.edu/~cs61c>

New School CS61C (1/3)

Personal
Mobile
Devices



New School CS61C (2/3)

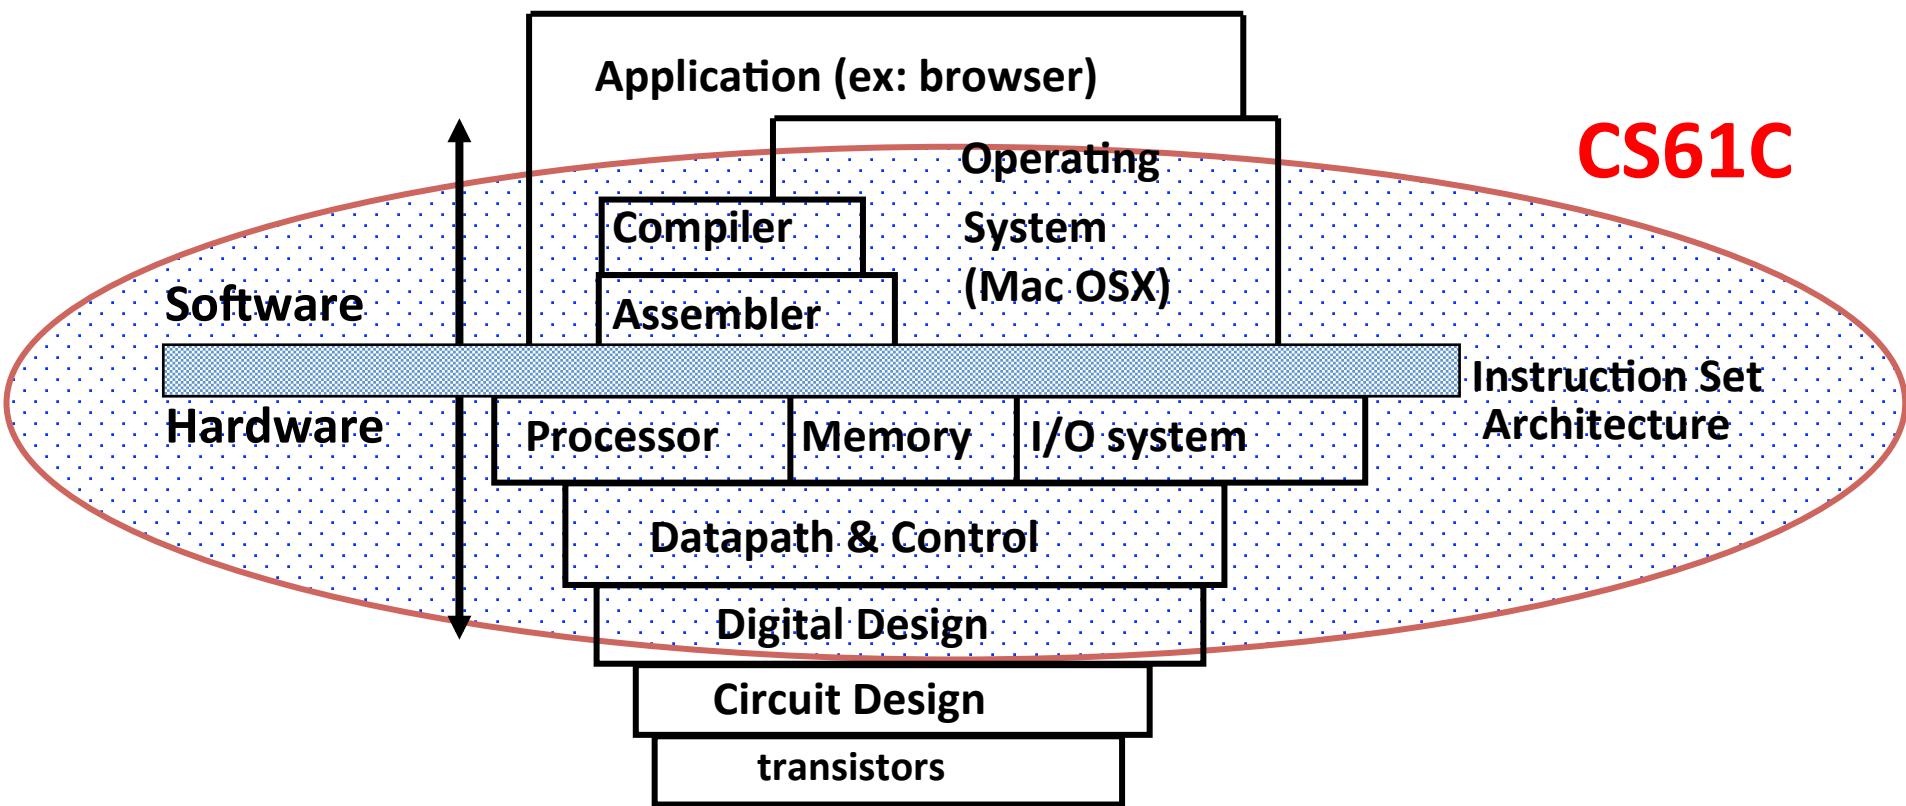


New School CS61C (3/3)

A perspective view looking down a long aisle in a data center. Both sides are filled with tall server racks, their front panels illuminated with a bright blue light. The floor is a polished grey, and the ceiling is high with industrial lighting fixtures. In the distance, a set of double doors is visible.

My other computer
is a data center

Old Machine Structures



New-School Machine Structures (It's a bit more complicated!)

Lab 13

Software

Hardware

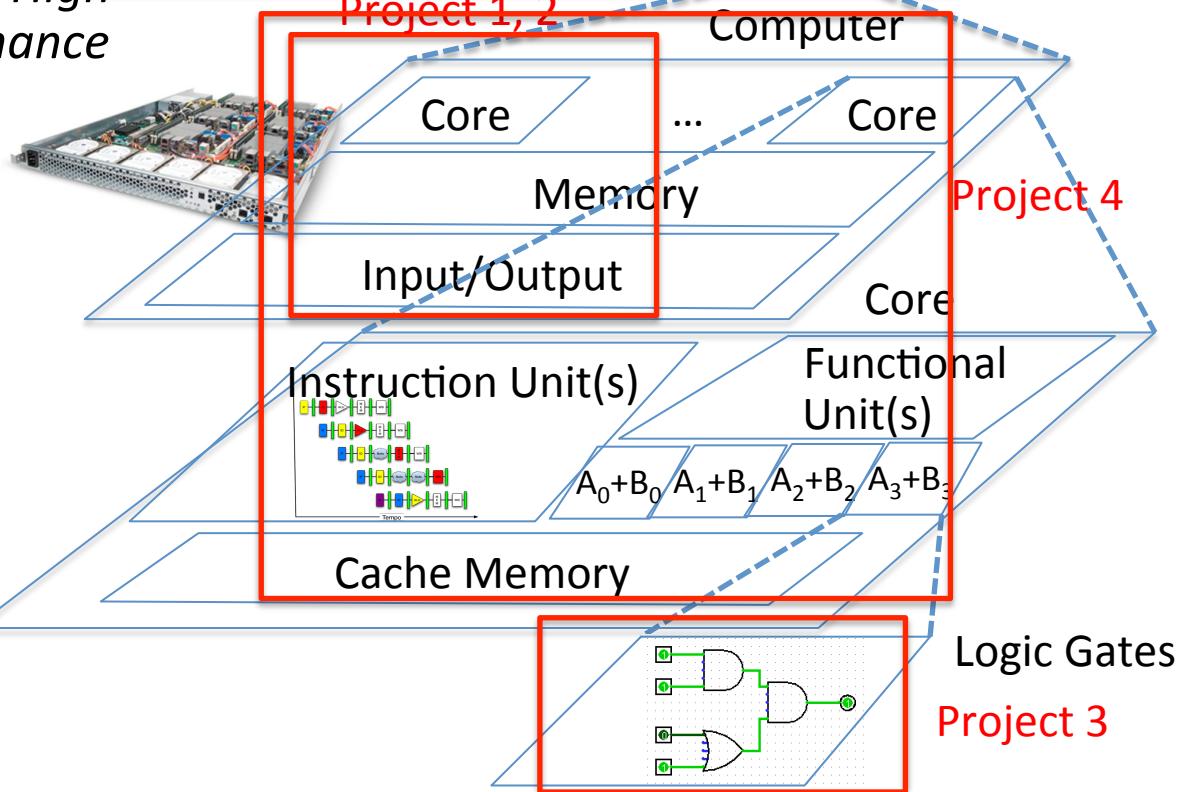
Leverage
Parallelism &
Achieve High
Performance

- Parallel Requests
Assigned to computer
e.g., Search “Katz”
- Parallel Threads
Assigned to core
e.g., Lookup, Ads
- Parallel Instructions
 >1 instruction @ one time
e.g., 5 pipelined instructions
- Parallel Data
 >1 data item @ one time
e.g., Add of 4 pairs of words
- Hardware descriptions
All gates functioning in parallel at same time
- Programming Languages

Warehouse Scale Computer



Smart Phone



CS61c is NOT about C Programming

- It's about the hardware-software interface
 - What does the programmer need to know to achieve the highest possible performance
- Languages like C are closer to the underlying hardware, unlike languages like Python!
 - Allows us to talk about key hardware features in higher level terms
 - Allows programmer to explicitly harness underlying hardware parallelism for high performance: “programming for performance”

Great Ideas in Computer Architecture

1. Design for Moore's Law
2. Abstraction to Simplify Design
3. Make the Common Case Fast
4. Dependability via Redundancy
5. Memory Hierarchy
6. Performance via Parallelism/Pipelining/
Prediction

Powers of Ten inspired 61C Overview

- Going Top Down cover 3 Views
 1. Architecture (when possible)
 2. Physical Implementation of that architecture
 3. Programming system for that architecture and implementation (when possible)
- See <http://www.powersof10.com/film>

Earth

10^7 meters

Powers of Ten™ (1977)
by EamesOffice

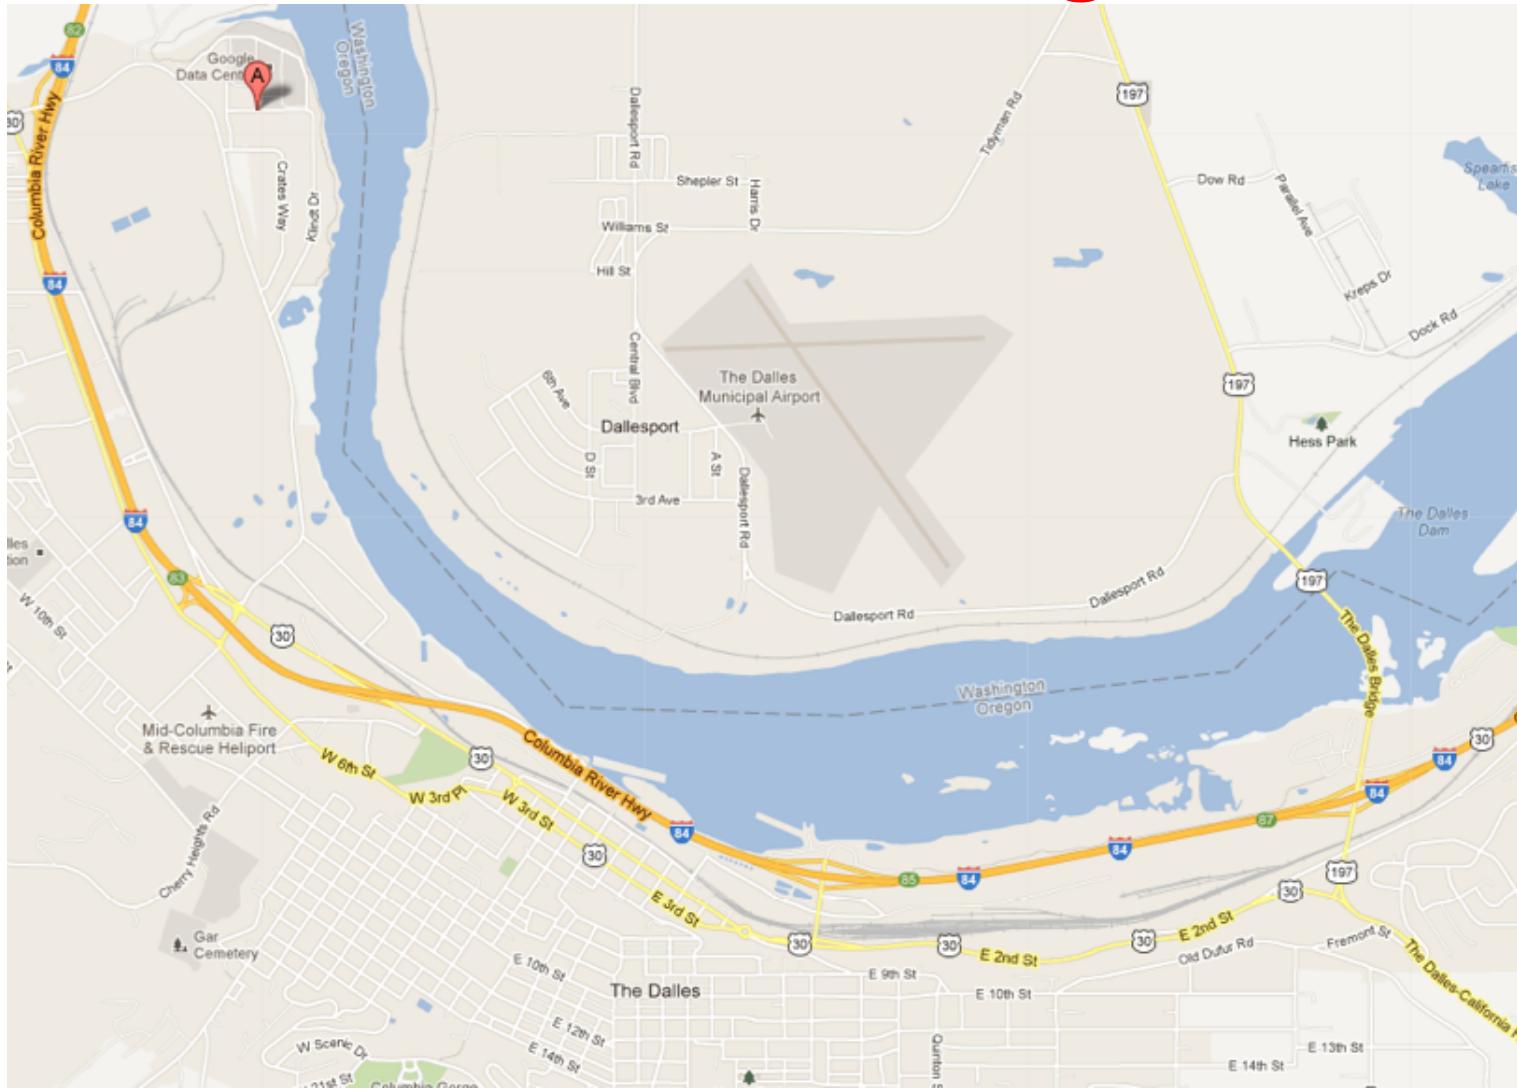
10 million meters

10^7
meters



The Dalles, Oregon

10^4 meters



The Dalles, Oregon 10^4 meters



Google's Oregon WSC 10^3 meters



10^4 meters

Google's Oregon WSC



10 kilometers

10^2 meters



10^3 meters



Google Warehouse

- 90 meters by 75 meters, 10 Megawatts
- Contains 40,000 servers, 190,000 disks
- Power Utilization Effectiveness: 1.23
 - 85% of 0.23 overhead goes to cooling losses
 - 15% of 0.23 overhead goes to power losses
- Contains 45, 40-foot long containers
 - 8 feet x 9.5 feet x 40 feet
- 30 stacked as double layer, 15 as single layer

10^2 meters

Containers in WSCs

100 meters



Google Container

10^1 meters

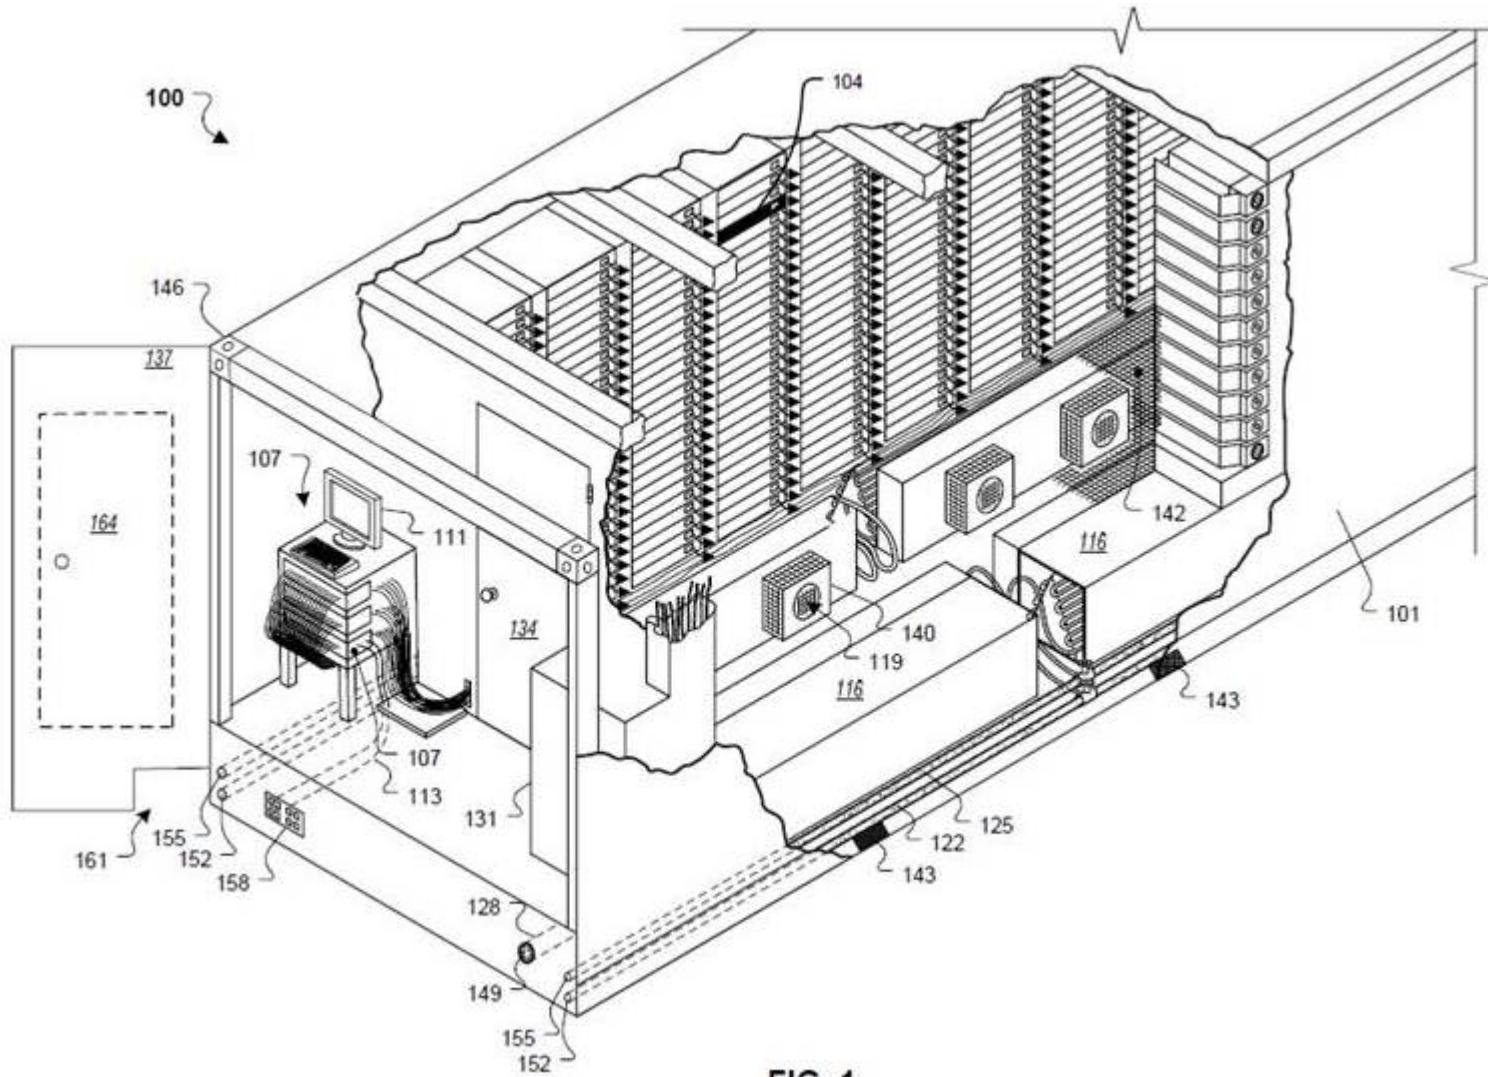
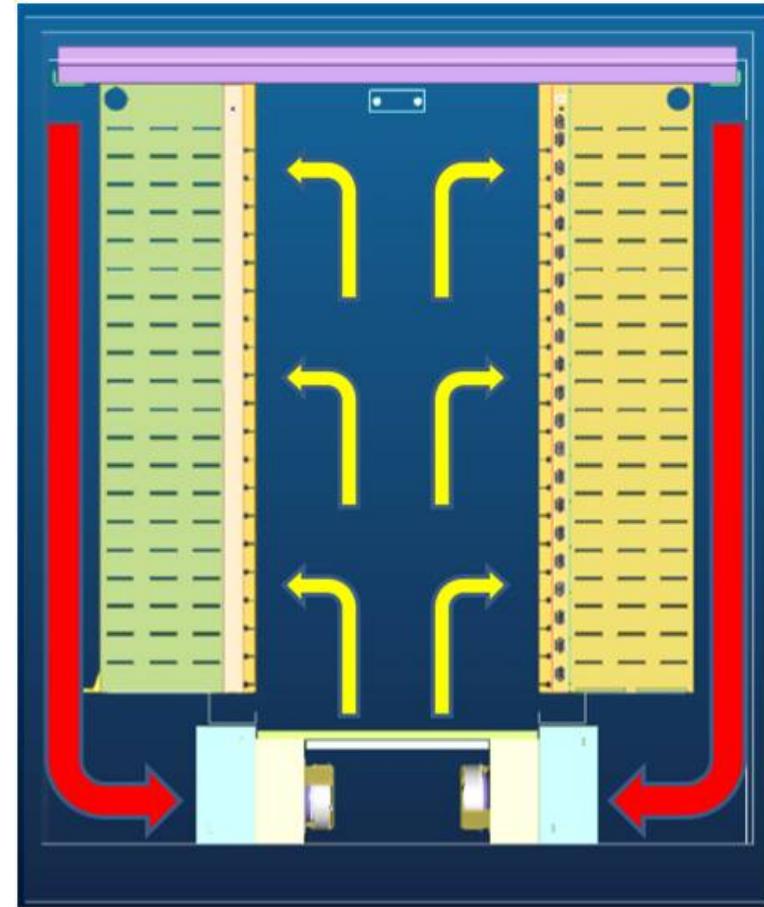


FIG. 1

Google Container

10^0 meters



- 2 long rows, each with 29 racks
- Cooling below raised floor
- Hot air returned behind racks

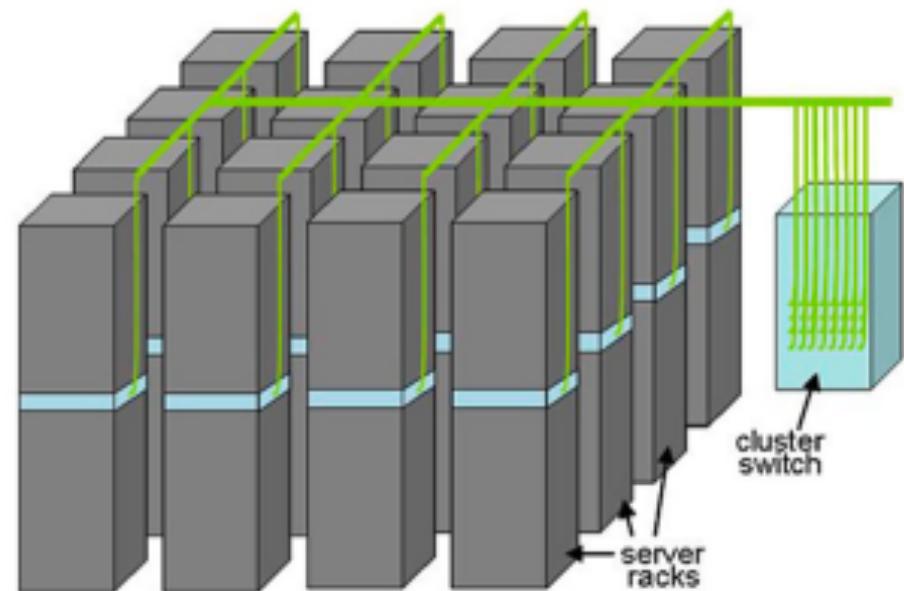
Equipment Inside a Container



Server (in rack format):



7 foot Rack: servers + Ethernet local area network switch in middle (“rack switch”)



Array (aka cluster):
server racks + larger local area network switch (“array switch”) 10X faster => cost 100X: cost $f(N^2)$

10^0 meters

Google Rack

- Google rack with 20 servers + Network Switch in the middle
- 48-port 1 Gigabit/sec Ethernet switch every other rack
- Array switches connect to racks via multiple 1 Gbit/s links
- 2 datacenter routers connect to array switches over 10 Gbit/s links

1 meter



Programming WSC: Word Count in Spark's Python API

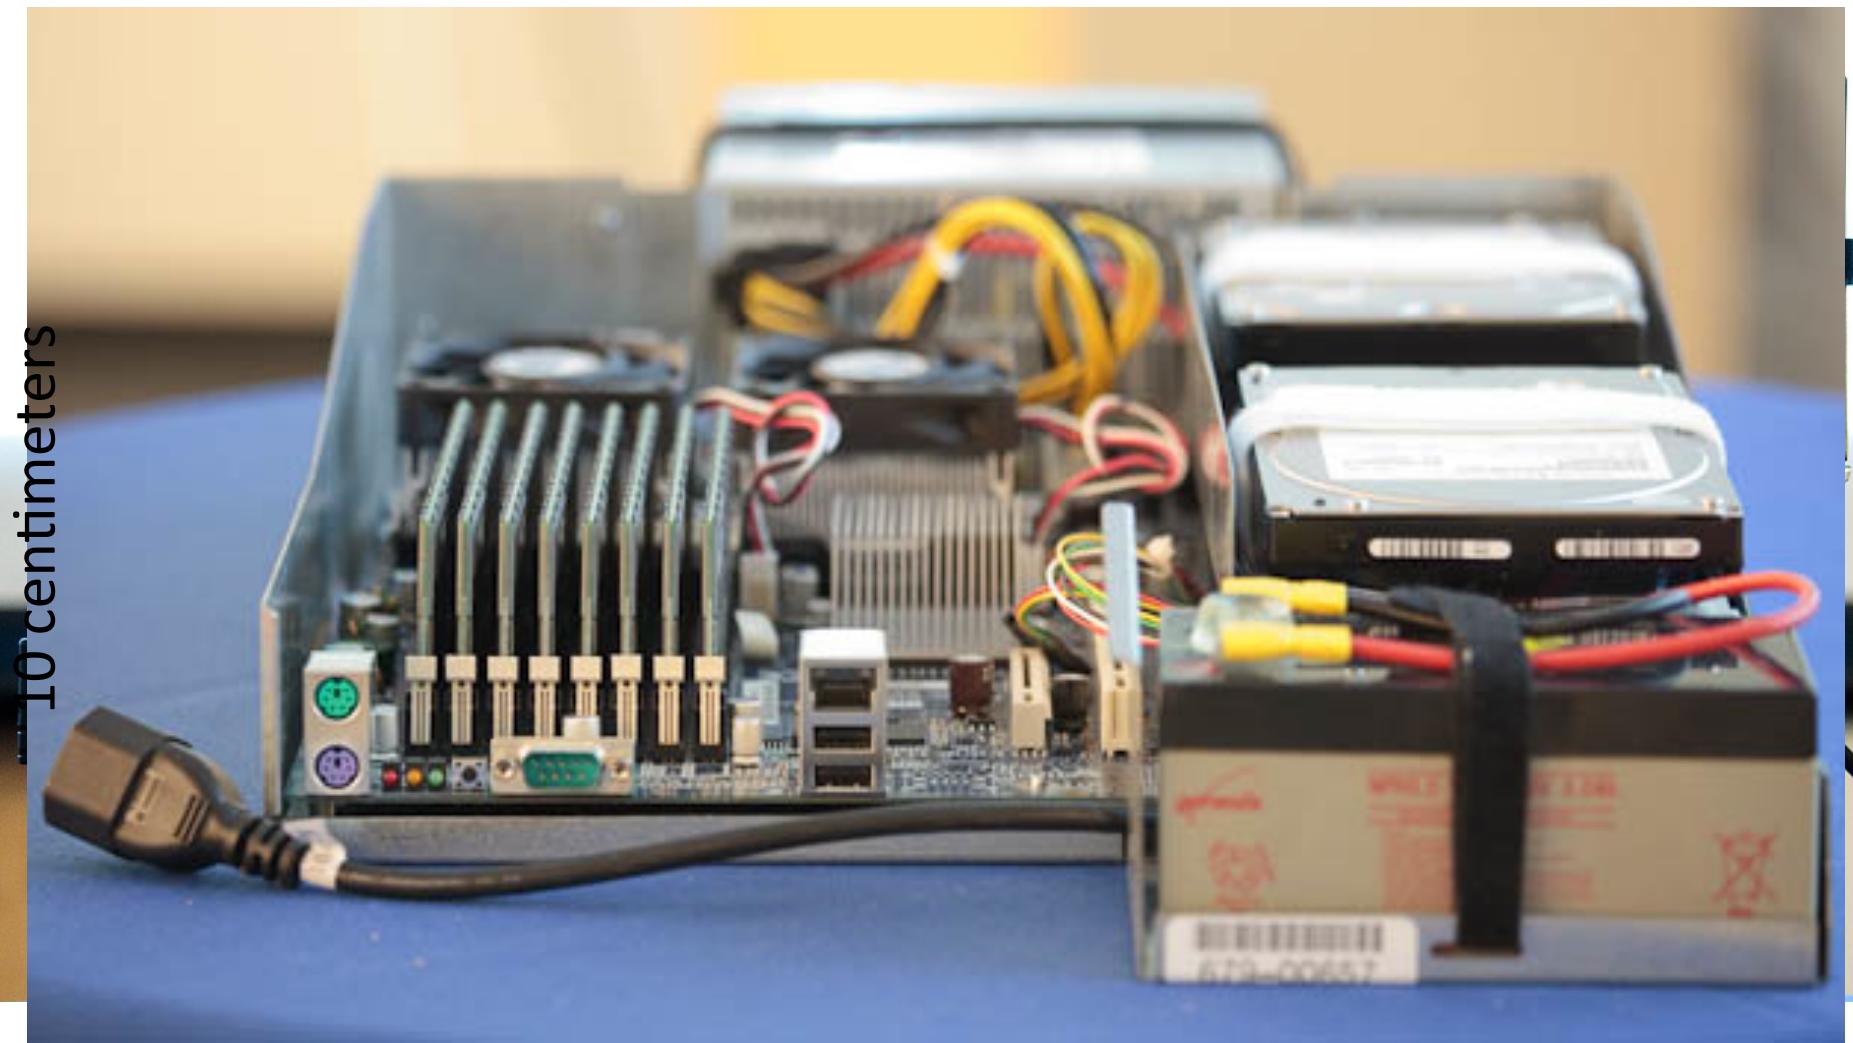
```
// RDD: primary abstraction of a distributed collection of items  
file = sc.textFile("hdfs://...")  
// Two kinds of operations:  
// Actions: RDD → Value  
// Transformations: RDD → RDD  
// e.g. flatMap, Map, reduceByKey  
file.flatMap(lambda line: line.split())  
    .map(lambda word: (word, 1))  
    .reduceByKey(lambda a, b: a + b)
```

Great Ideas in Computer Architecture

1. *Design for Moore's Law*
 - *WSC, Container, Rack*
2. Abstraction to Simplify Design
3. Make the Common Case Fast
4. *Dependability via Redundancy*
 - *Multiple WSCs, Multiple Racks, Multiple Switches*
5. Memory Hierarchy
6. *Performance via Parallelism/Pipelining/
Prediction*
 - *Task level Parallelism, Data Level Parallelism*

Google Server Internals

10^{-1} meters



Google Board Details

- Supplies only 12 volts
- Battery per board vs. large battery room
 - Improves PUE: 99.99% efficient local battery vs 94% for battery room
- 2 SATA Disk Drives
 - 1 Terabyte capacity each
 - 3.5 inch disk drive
 - 7200 RPM
- 2 AMD Opteron Microprocessors
 - Dual Core, 2.2 GHz
- 8 DIMMs
 - 8 GB DDR2 DRAM
- 1 Gbit/sec Ethernet Network Interface Card

Programming Multicore Microprocessor: OpenMP

```
#include <omp.h>
#include <stdio.h>
static long num_steps = 100000;
int value[num_steps];
int reduce()
{
    int i;    int sum = 0;
#pragma omp parallel for private(x) reduction(+:sum)
    for (i=1; i<= num_steps; i++){
        sum = sum + value[i];
    }
}
```

Great Ideas in Computer Architecture

1. *Design for Moore's Law*
 - More transistors = Multicore + SIMD
2. Abstraction to Simplify Design
3. Make the Common Case Fast
4. Dependability via Redundancy
5. *Memory Hierarchy*
 - More transistors = Cache Memories
6. *Performance via Parallelism/Pipelining/ Prediction*
 - Thread-level Parallelism

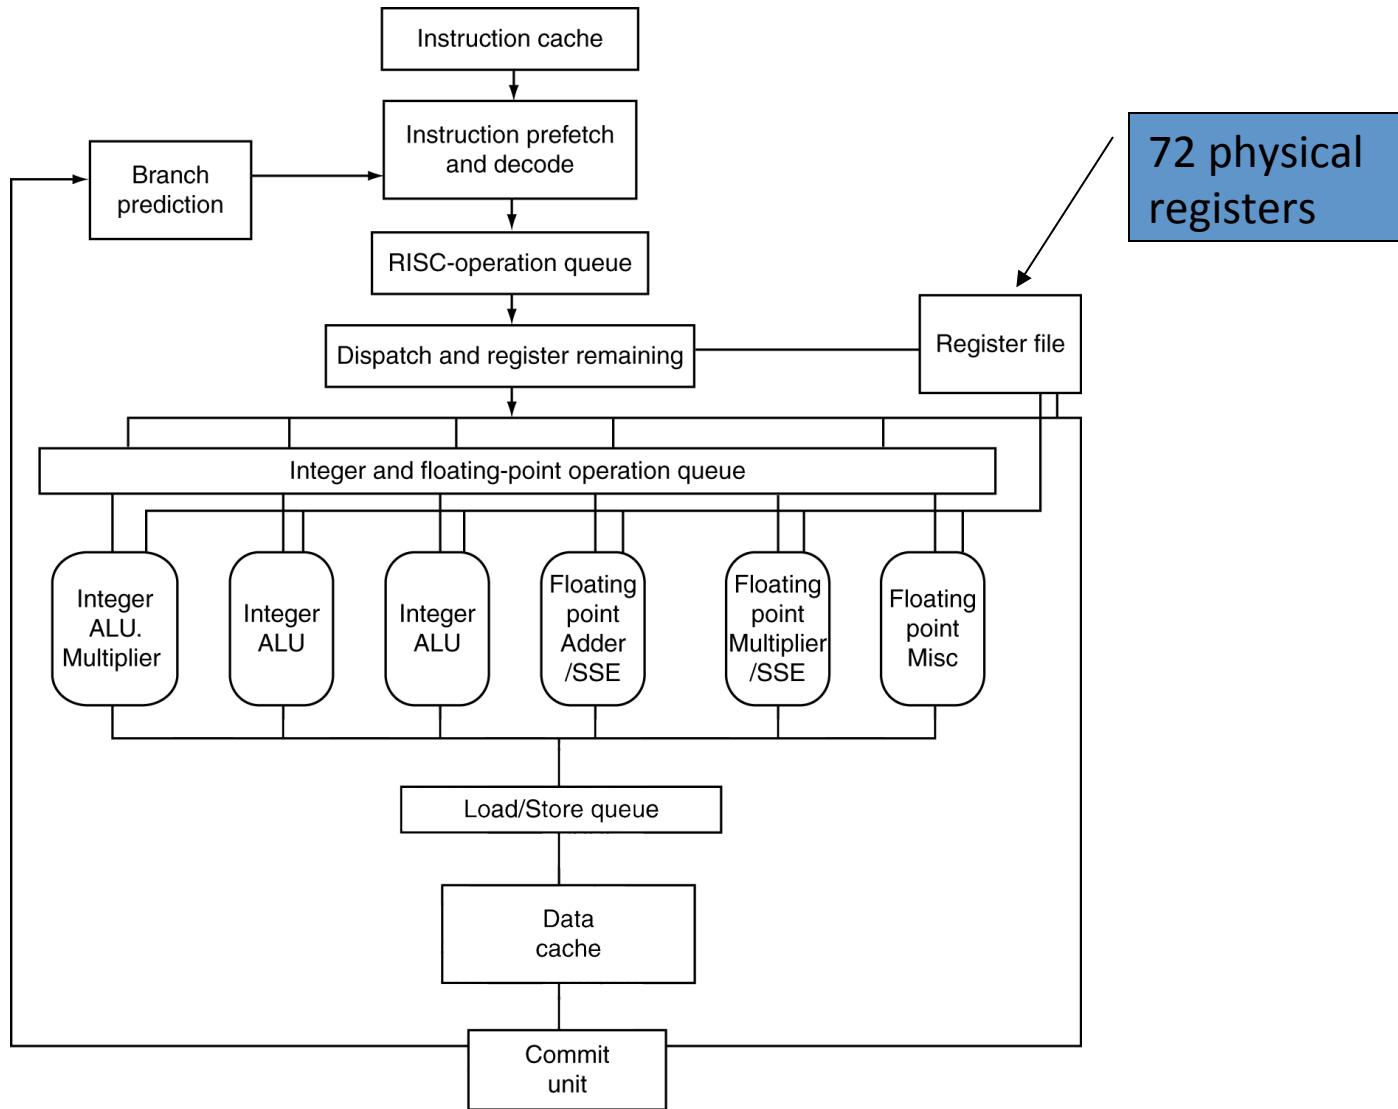
10^{-2} meters

AMD Opteron Microprocessor

centimeters

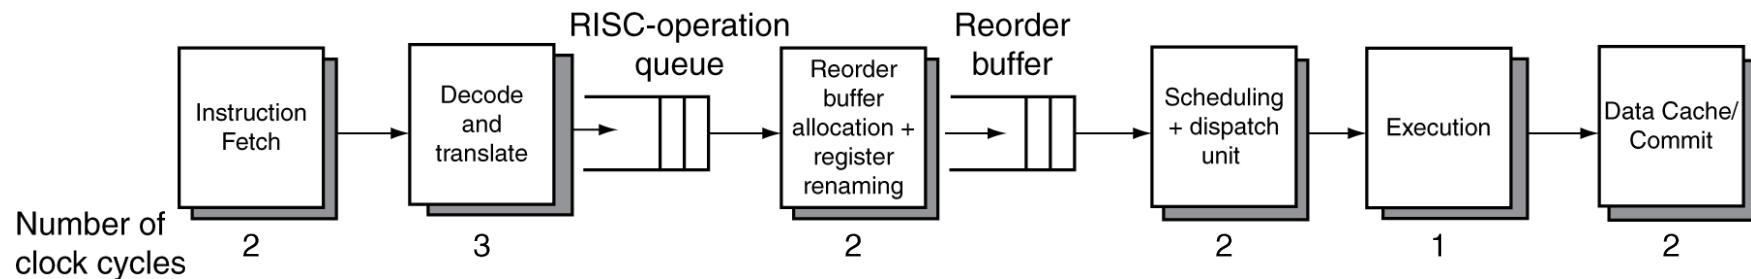


AMD Opteron Microarchitecture



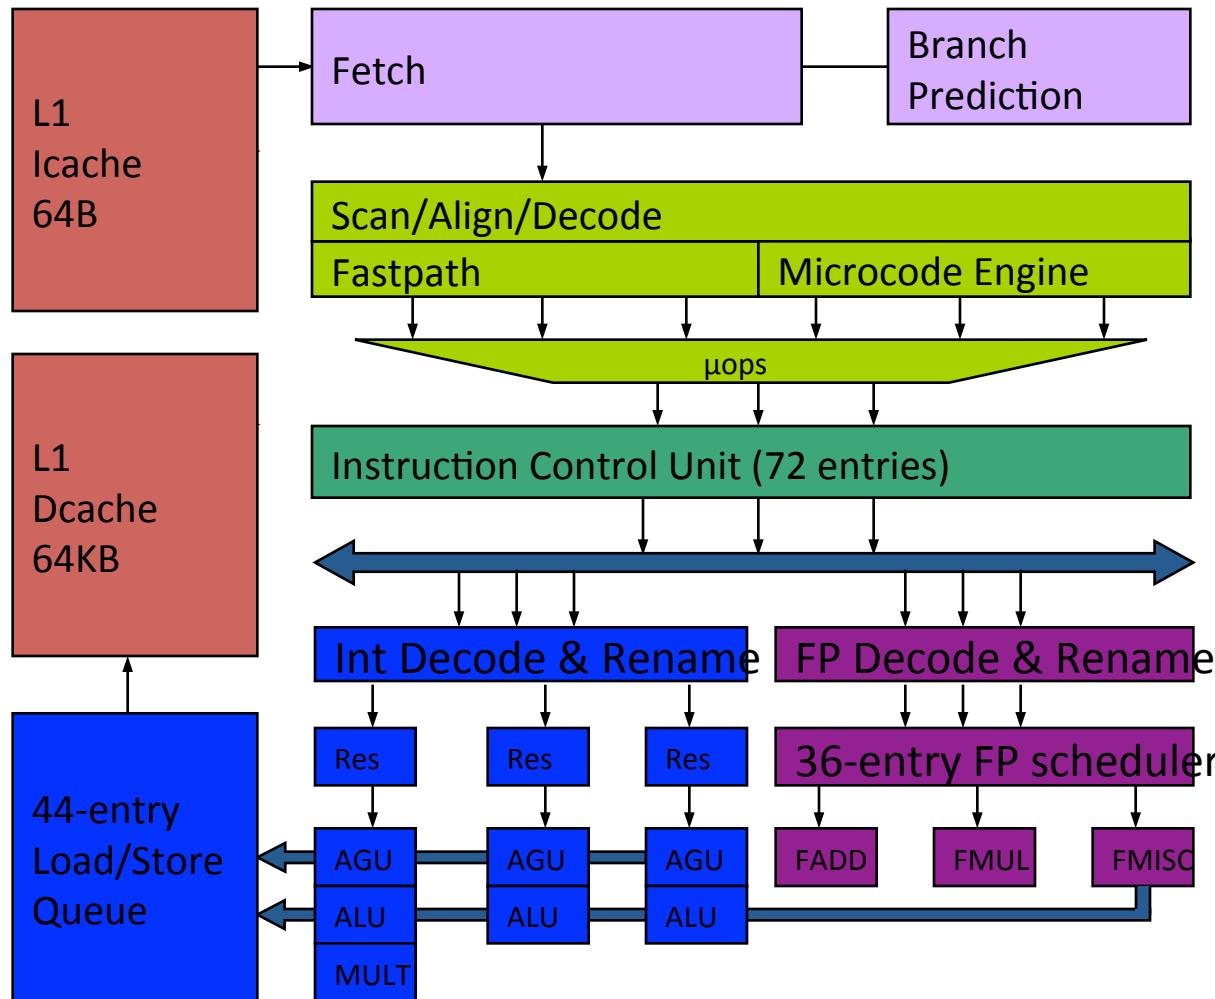
AMD Opteron Pipeline Flow

- For integer operations



- 12 stages (Floating Point is 17 stages)
- Up to 106 RISC-ops in progress

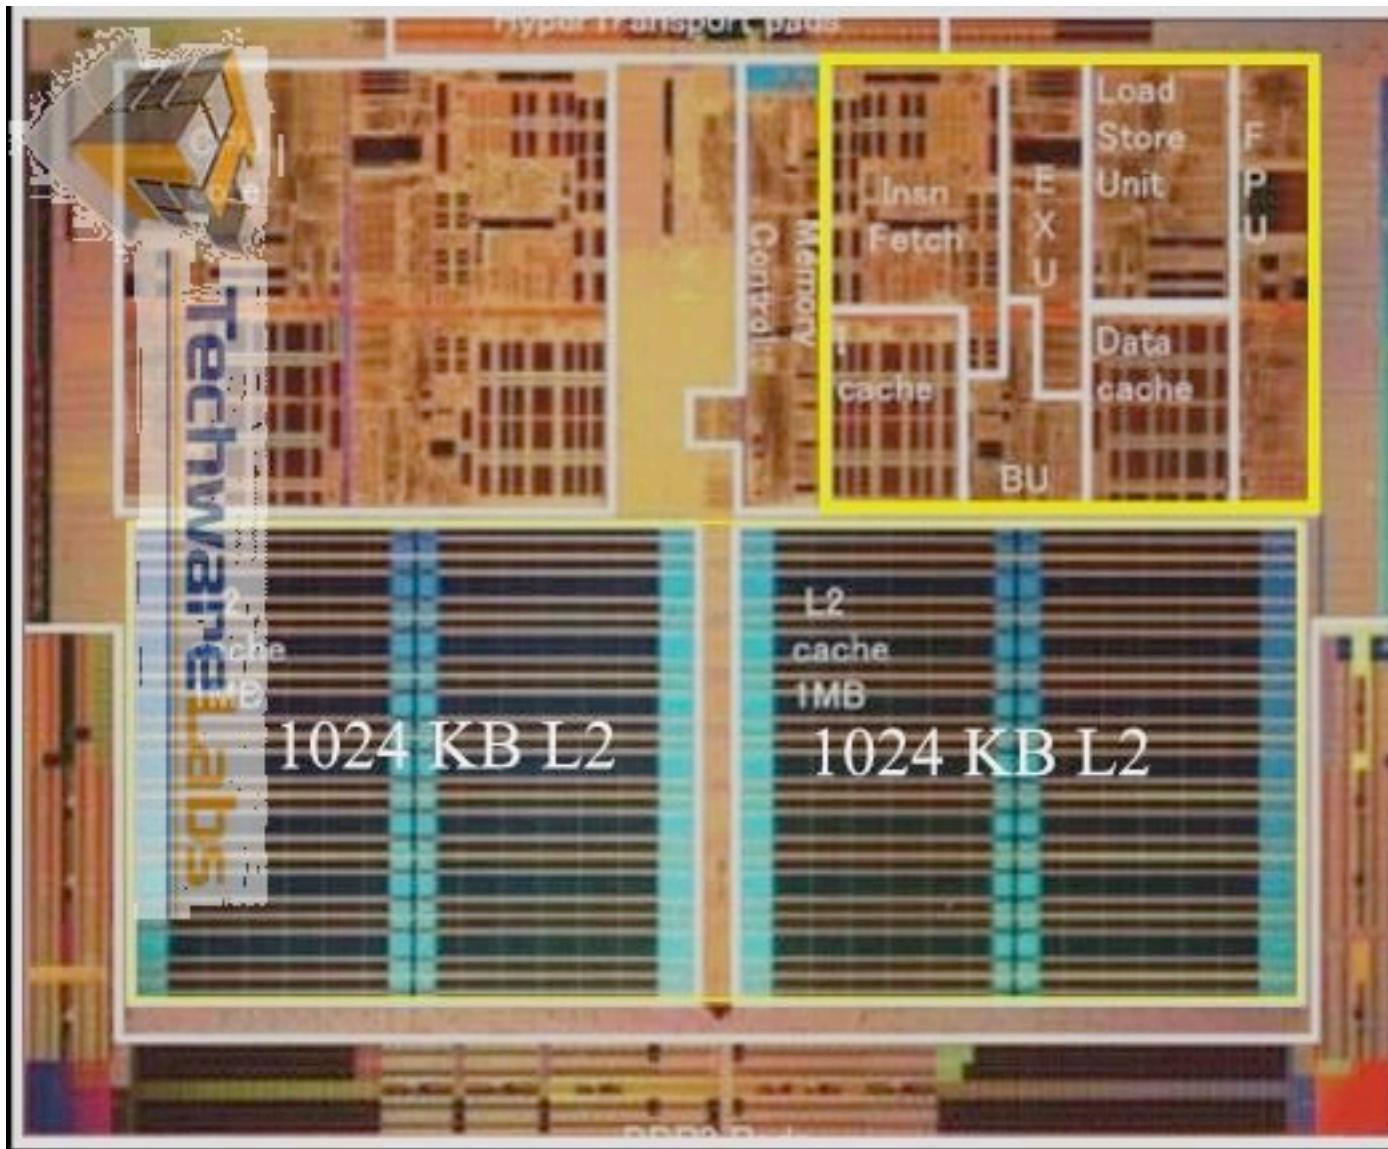
AMD Opteron Block Diagram



10^{-2} meters

AMD Opteron Microprocessor

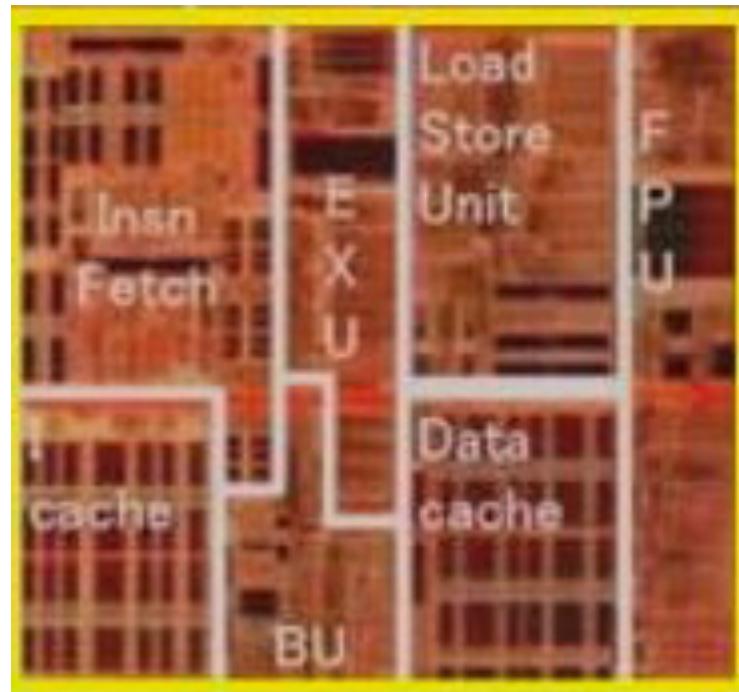
centimeters



10^{-3} meters

AMD Opteron Core

millimeters



Programming One Core: C with Intrinsics

```
void mmult(int n, float *A, float *B, float *C)
{
    for ( int i = 0; i < n; i+=4 )
        for ( int j = 0; j < n; j++ )
    {
        __m128 c0 = _mm_load_ps(C+i+j*n);
        for( int k = 0; k < n; k++ )
            c0 = _mm_add_ps(c0, _mm_mul_ps(_mm_load_ps(A+i+k*n),
                                              _mm_load1_ps(B+k+j*n)));
        _mm_store_ps(C+i+j*n, c0);
    }
}
```

Inner loop from gcc -O -S

Assembly snippet from innermost loop:

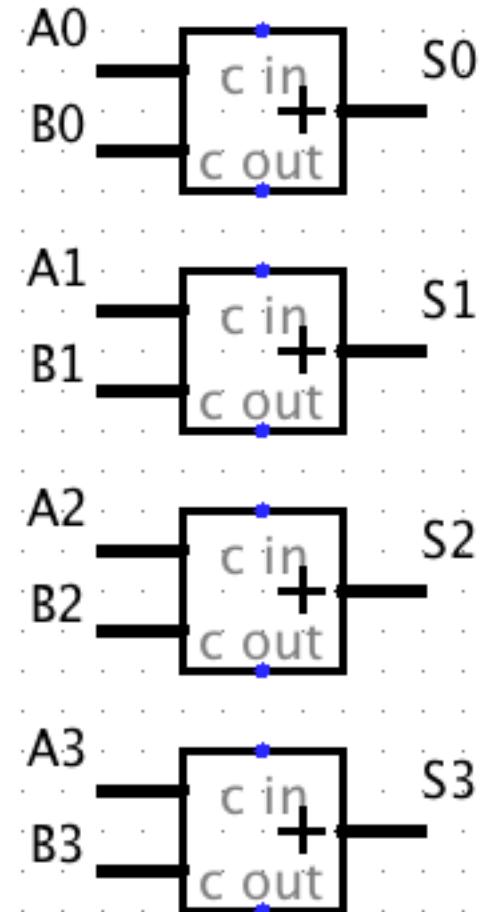
```
    movaps (%rax), %xmm9
    mulps %xmm0, %xmm9
    addps %xmm9, %xmm8
    movaps 16(%rax), %xmm9
    mulps %xmm0, %xmm9
    addps %xmm9, %xmm7
    movaps 32(%rax), %xmm9
    mulps %xmm0, %xmm9
    addps %xmm9, %xmm6
    movaps 48(%rax), %xmm9
    mulps %xmm0, %xmm9
    addps %xmm9, %xmm5
```

Great Ideas in Computer Architecture

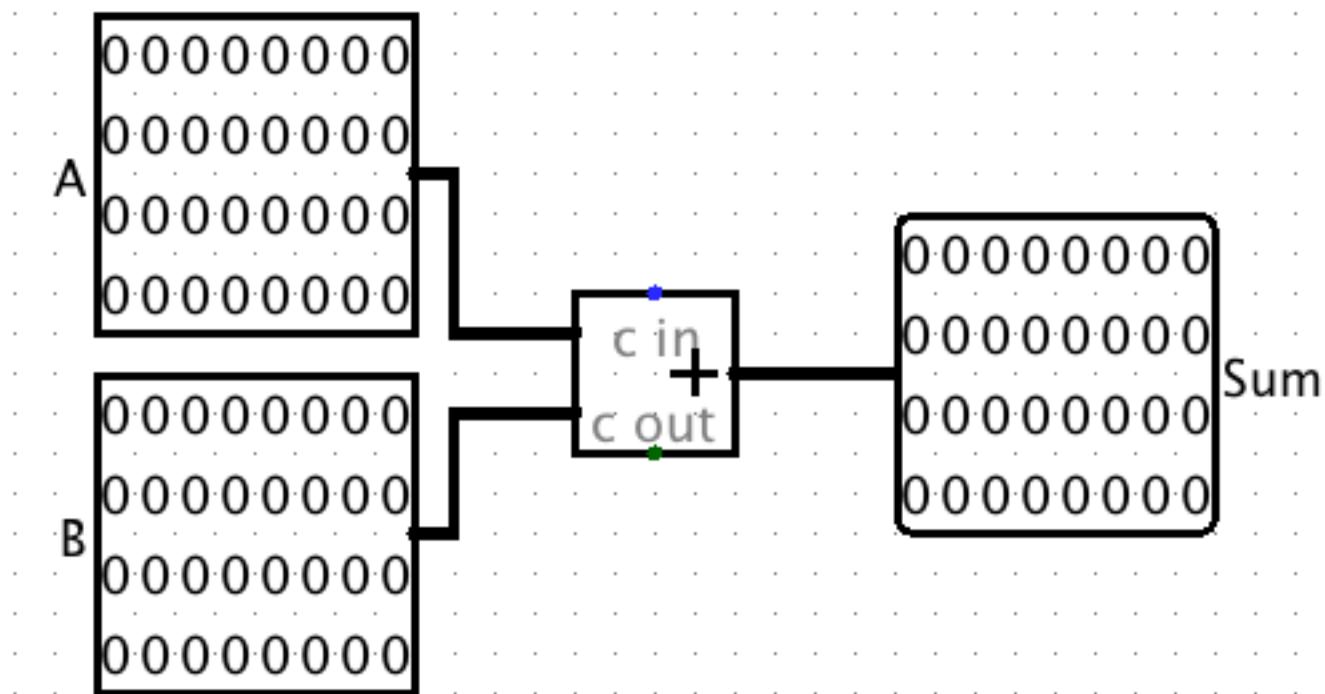
1. Design for Moore's Law
2. *Abstraction to Simplify Design*
 - *Instruction Set Architecture, Micro-operations*
3. Make the Common Case Fast
4. Dependability via Redundancy
5. Memory Hierarchy
6. *Performance via Parallelism/Pipelining/ Prediction*
 - *Instruction-level Parallelism (superscalar, pipelining)*
 - *Data-level Parallelism*

SIMD Adder

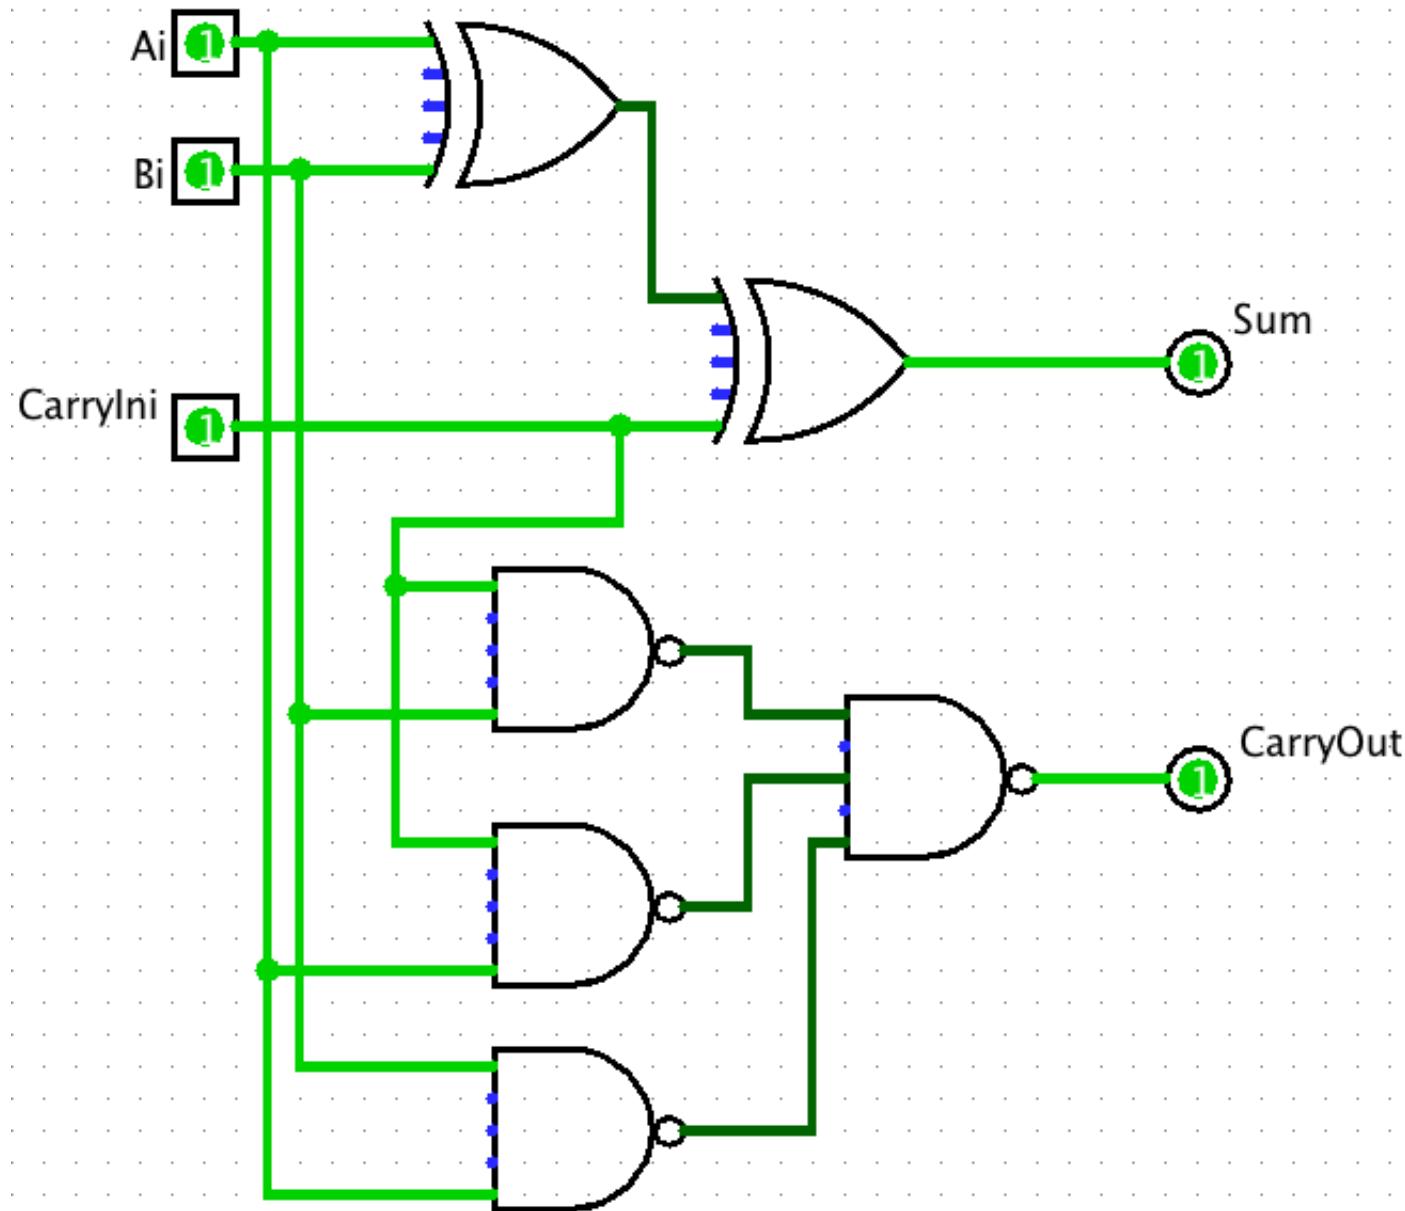
- Four 32-bit adders that operate in parallel
 - Data Level Parallelism



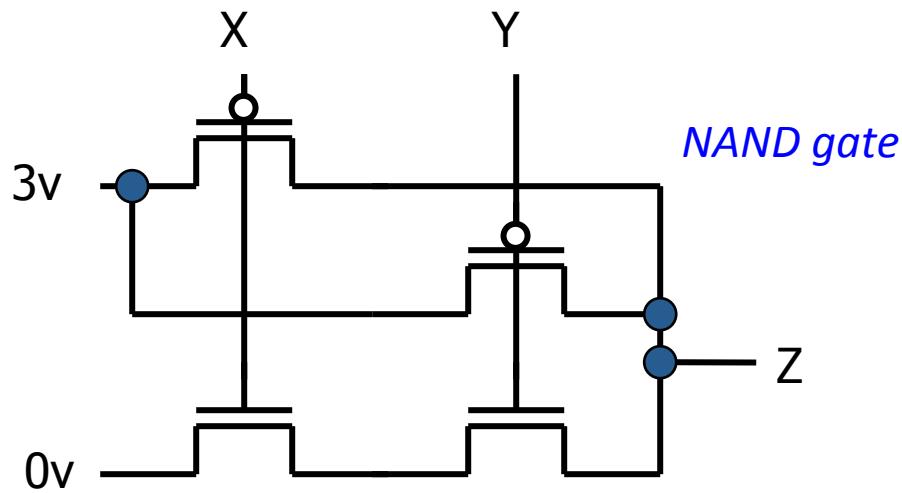
One 32-bit Adder



1 bit of 32-bit Adder



Complementary MOS Transistors (NMOS and PMOS) of NAND Gate

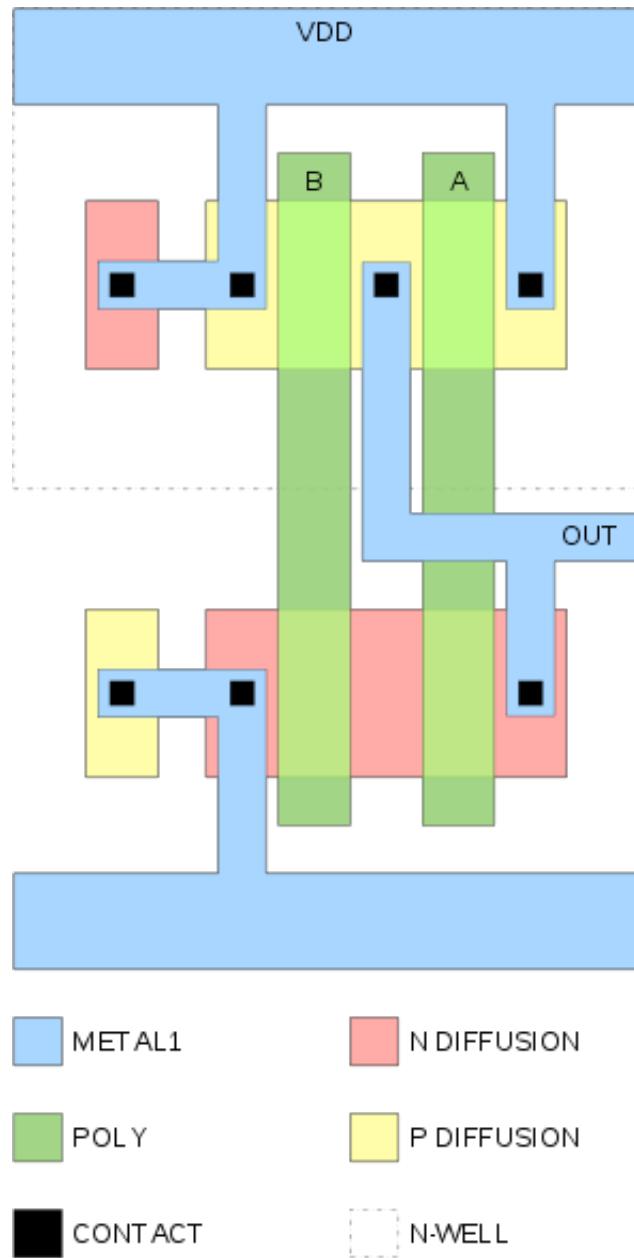


x	y	z
0 volts	0 volts	3 volts
0 volts	3 volts	3 volts
3 volts	0 volts	3 volts
3 volts	3 volts	0 volts

10^{-7} meters

Physical Layout of NAND Gate

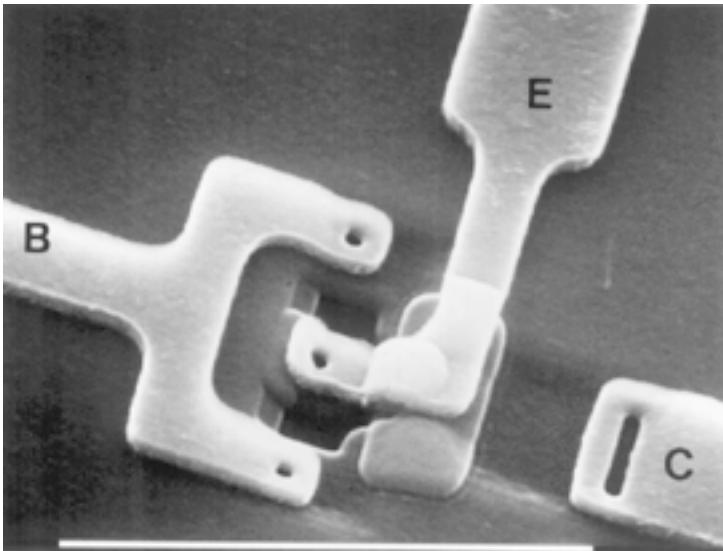
100 nanometers



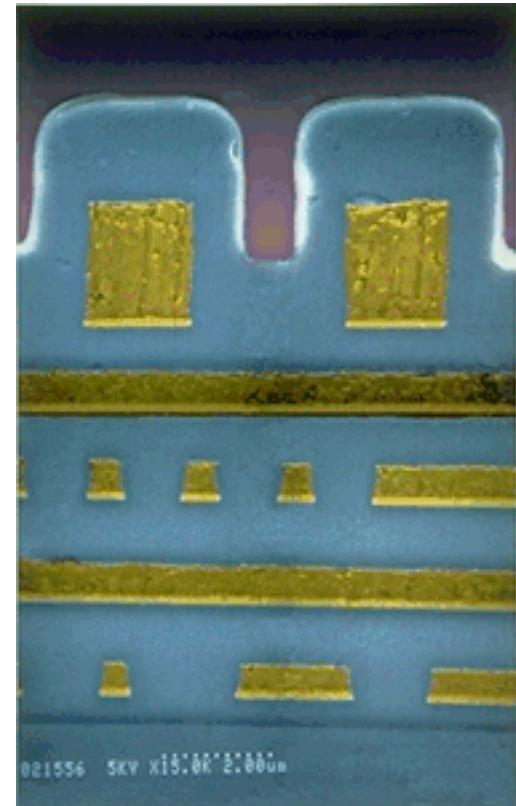
10^{-7} meters

Scanning Electron Microscope

100 nanometers



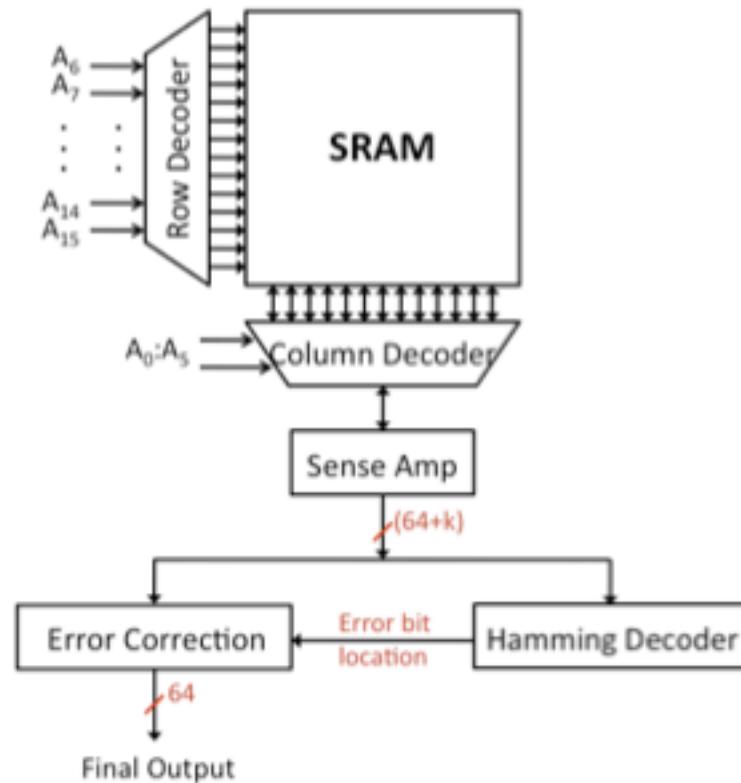
Top View



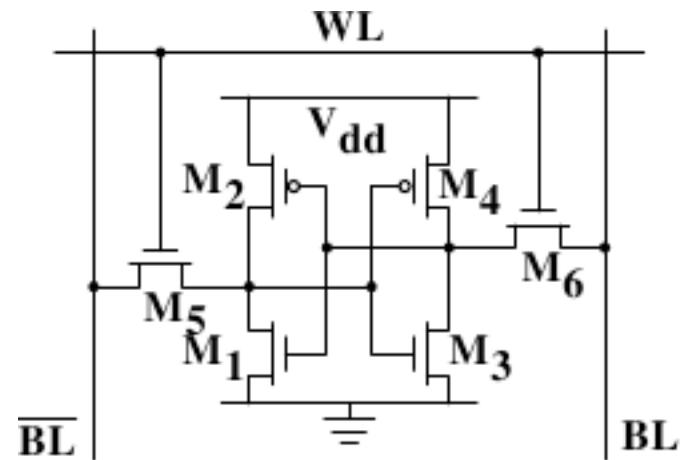
Cross Section

10^{-6} meters

Block Diagram of Static RAM



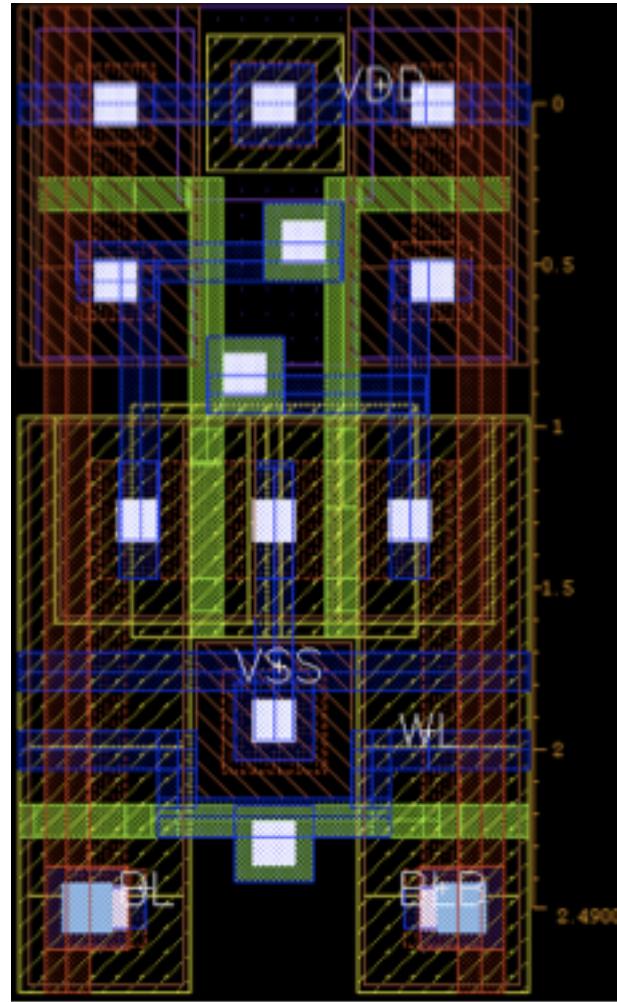
1 Bit SRAM in 6 Transistors



10^{-7} meters

Physical Layout of SRAM Bit

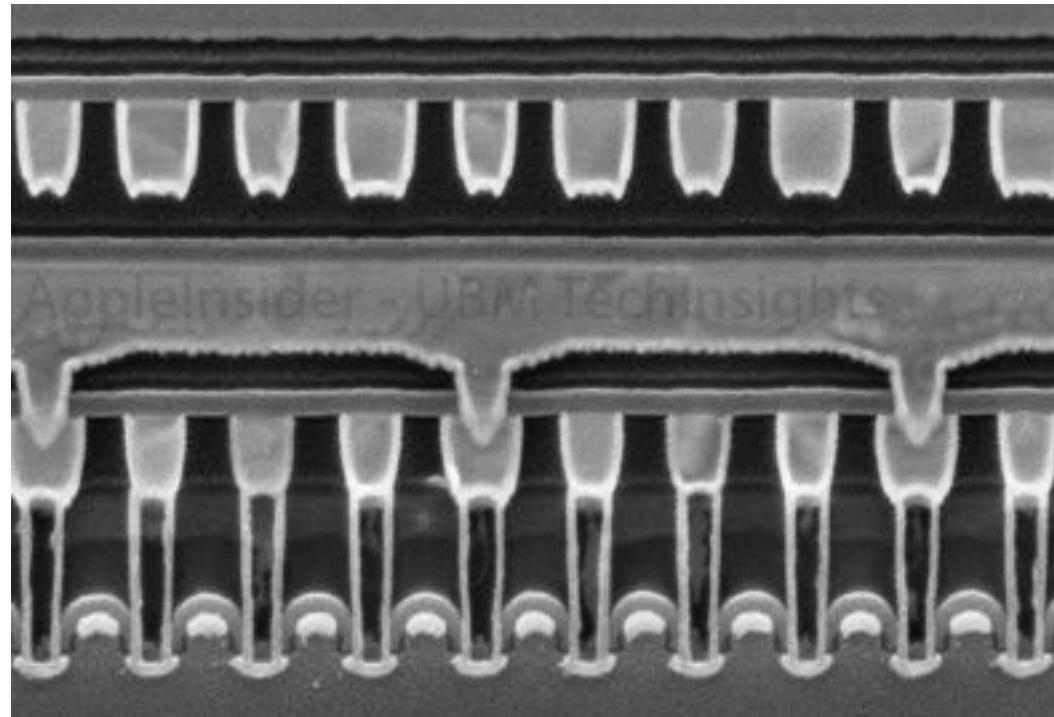
100 nanometers



10^{-7} meters

SRAM Cross Section

100 nanometers



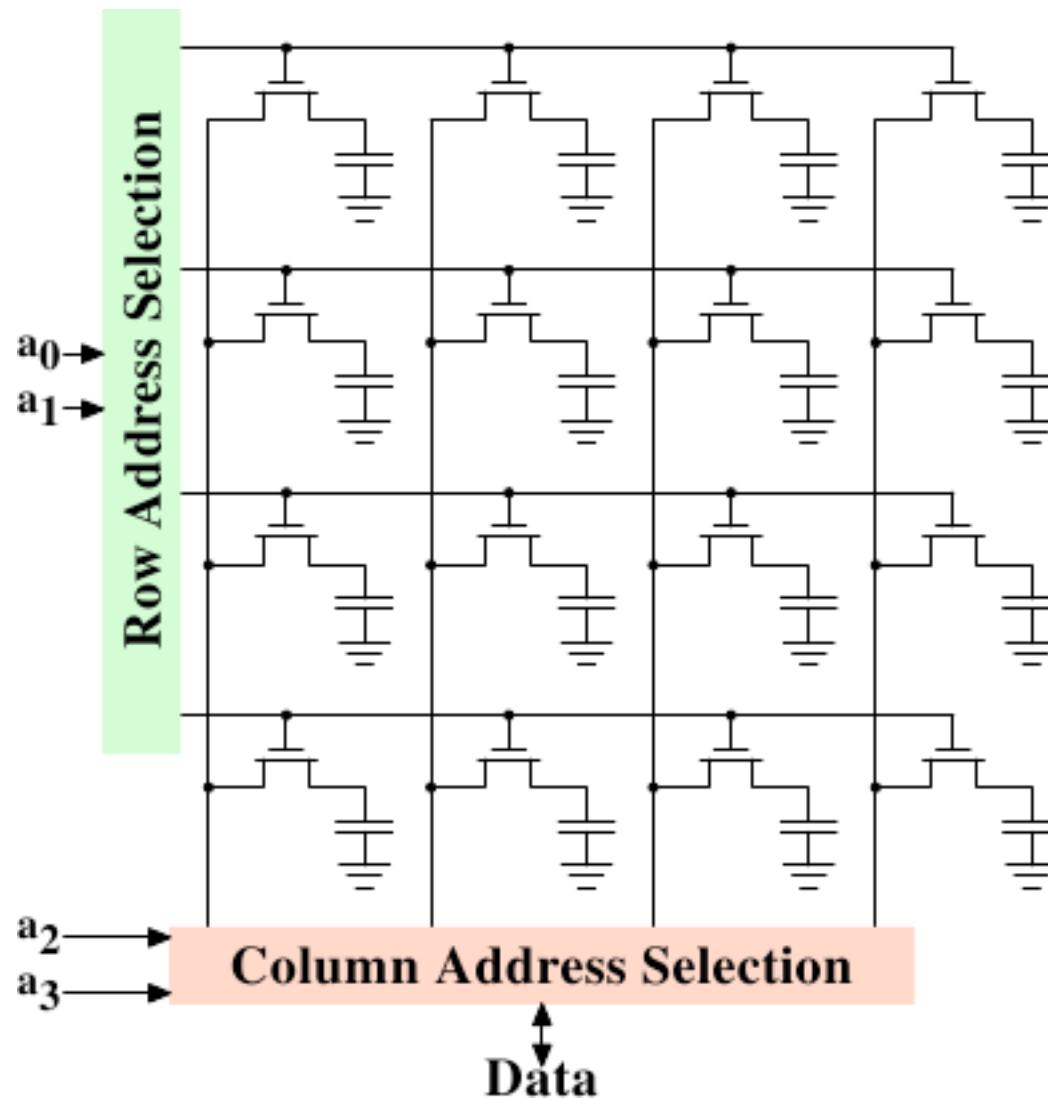
DIMM Module

- DDR = Double Data Rate
 - Transfers bits on Falling AND Rising Clock Edge
- Has Single Error Correcting, Double Error Detecting Redundancy (SEC/DED)
 - 72 bits to store 64 bits of data
 - Uses “Chip kill” organization so that if single DRAM chip fails can still detect failure
- Average server has 22,000 correctable errors and 1 uncorrectable error per year

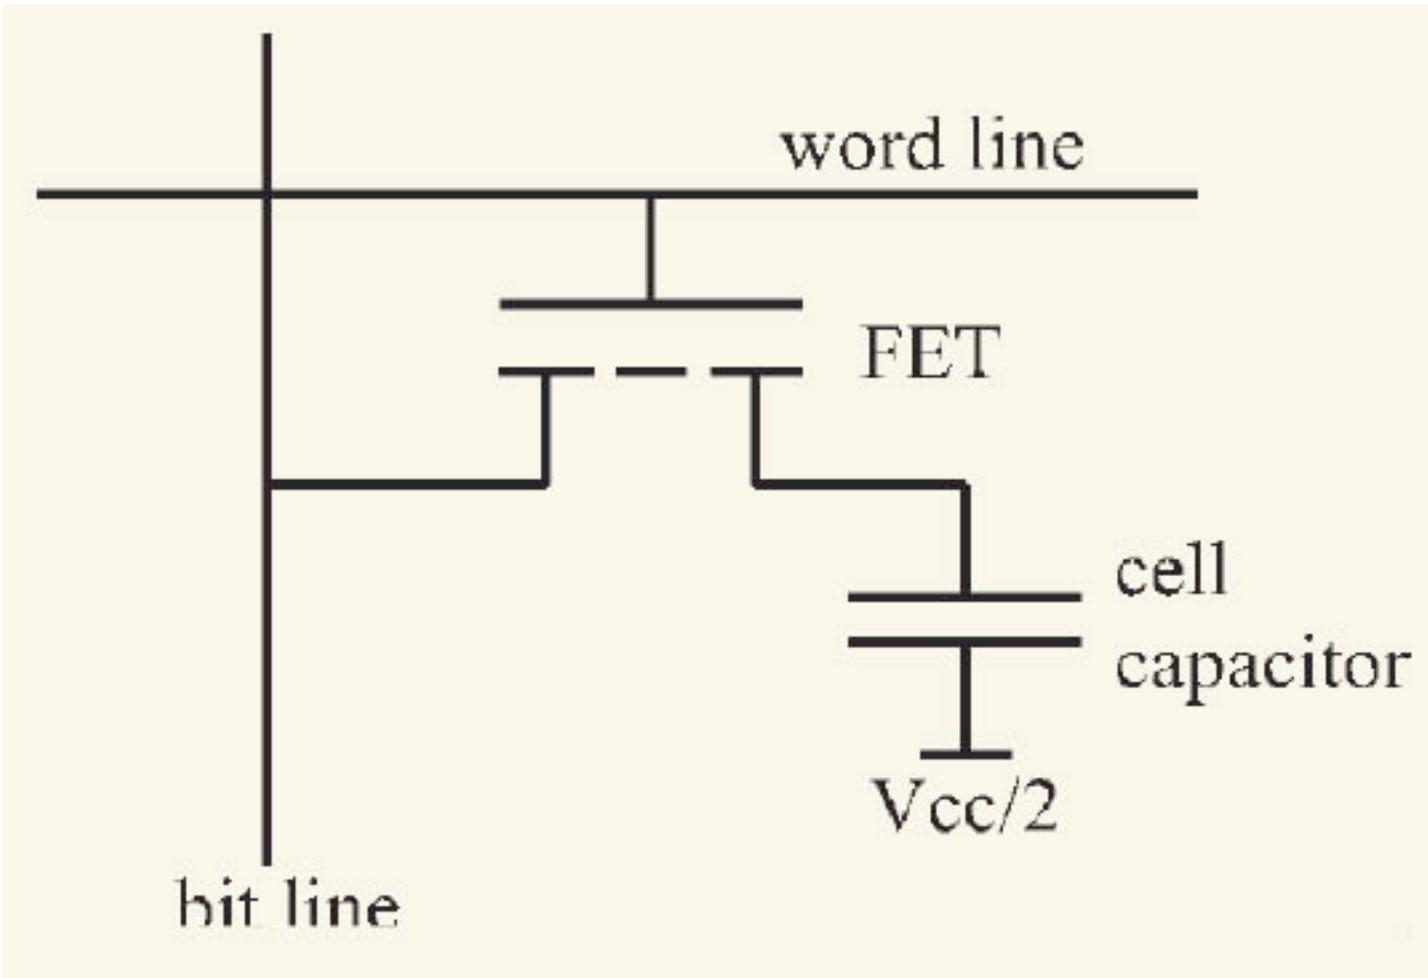
10^{-6} meters

DRAM Bits

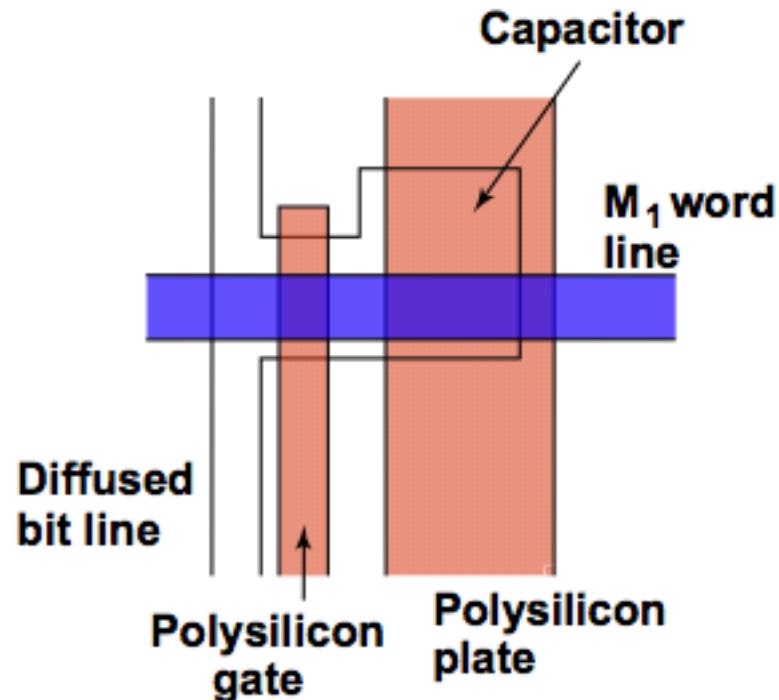
1 micron



DRAM Cell in Transistors



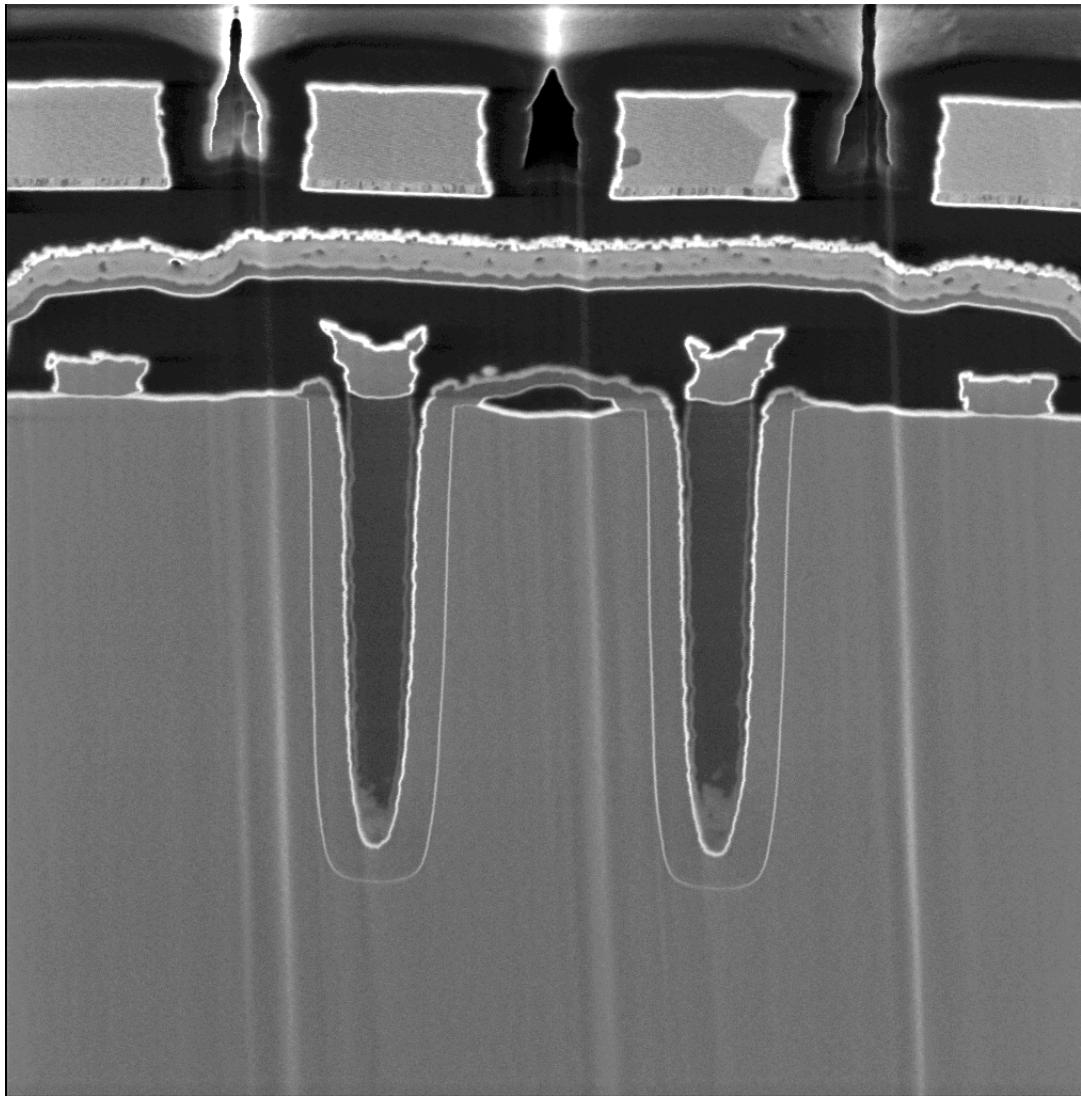
Physical Layout of DRAM Bit



10^{-7} meters

Cross Section of DRAM Bits

100 nanometers



AMD Dependability

- L1 cache data is SEC/DED protected
- L2 cache and tags are SEC/DED protected
- DRAM is SEC/DED protected with chipkill
- On-chip and off-chip ECC protected arrays include autonomous, background hardware scrubbers
- Remaining arrays are parity protected
 - Instruction cache, tags and TLBs
 - Data tags and TLBs
 - Generally read only data that can be recovered from lower levels

Programming Memory Hierarchy: Cache Blocked Algorithm

- The blocked version of the i-j-k algorithm is written simply as (A,B,C are submatrices of a, b, c)

```
for (i=0 ; i<N/r ; i++)
    for (j=0 ; j<N/r ; j++)
        for (k=0 ; k<N/r ; k++)
            C[i][j] += A[i][k]*B[k][j]
```

- r = block (sub-matrix) size (Assume r divides N)
- $X[i][j]$ = a sub-matrix of X , defined by block row i and block column j

Great Ideas in Computer Architecture

1. *Design for Moore's Law*
 - *Higher capacities caches and DRAM*
2. Abstraction to Simplify Design
3. Make the Common Case Fast
4. *Dependability via Redundancy*
 - *Parity, SEC/DEC*
5. *Memory Hierarchy*
 - *Caches, TLBs*
6. *Performance via Parallelism/Pipelining/Prediction*
 - *Data-level Parallelism*

Course Summary

- As the field changes, cs61c had to change too!
- It is still about the software-hardware interface
 - Programming for performance!
 - Parallelism: Task-, Thread-, Instruction-, and Data-MapReduce, OpenMP, C, SSE intrinsics
 - Understanding the memory hierarchy and its impact on application performance
- Interviewers ask what you did this semester!

Administrivia

- No more assignments!
- Upcoming Lecture Schedule
 - 8/11: Summary, What's Next? (+ HKN reviews)
 - Project 4 Competition Winners Announced
 - (Today)
 - 8/12: No Lecture, I'll have OH in this room

Administrivia

- Final Exam is this Thursday (8/13)
 - 9am-12pm, 10 Evans
 - Three 8.5"x11", double-sided, handwritten cheat-sheets
 - Broken into MT1 section, MT2 section, and post-MT2 section (so that the clobber can be applied)
 - Review slides posted on Piazza
- Extra OH - see Piazza post
- Labs today are final checkoffs + OH
- Discussions tomorrow are OH

Competition Prize Presentation

What Next? Classes

Welcome to the upper-div!

- CS150/EE141 (fall) if you liked digital systems design or transistor-level design
- CS152 (spring) if you liked computer architecture
- CS162 (operating systems and system programming) for more low-level software
- CS164 for C.A.L.L.
- CS161 for security

What Next?

Architecture research at Berkeley:

- <https://eecs.berkeley.edu/Research/Areas/ARC/>
- Build real chips!

Become a CS61C Staff Member!

- General path:
 - Lab Assistant to Reader or Tutor to TA
- Lab Assistants earn credit
- Readers/Tutors are paid
- TAs are paid + tuition + units
- CS61C Fa15: Currently 668 enrolled

Thanks to all Staff!

TAs:

- Nathaniel Mailoa (Head TA)
- Jay Patel (Head TA)
- Derek Ahmed
- Rebecca Herman
- Harrison Wang
- Jeffrey Wettstein

Tutors:

- Brenton Chu
- Nicolas Stone
- Alex Sung
- Austin Tai
- Michelle Tsai

Readers:

- Dasheng Chen
- Molly Zhai